

Distributed Systems

DAT520 - Spring 2026

Chapter 1 - Introduction

Prof. Hein Meling

Programming Abstractions

Encapsulate distributed interactions within

- **Abstractions** with *well-defined interfaces*
- To help us:
 - Reason about the **correctness** of distributed applications

Well-known Programming Abstractions

- Sequential programming abstractions
 - Set, Map, List
- Concurrent programming abstractions
 - Thread, Semaphore, Monitor, Locks

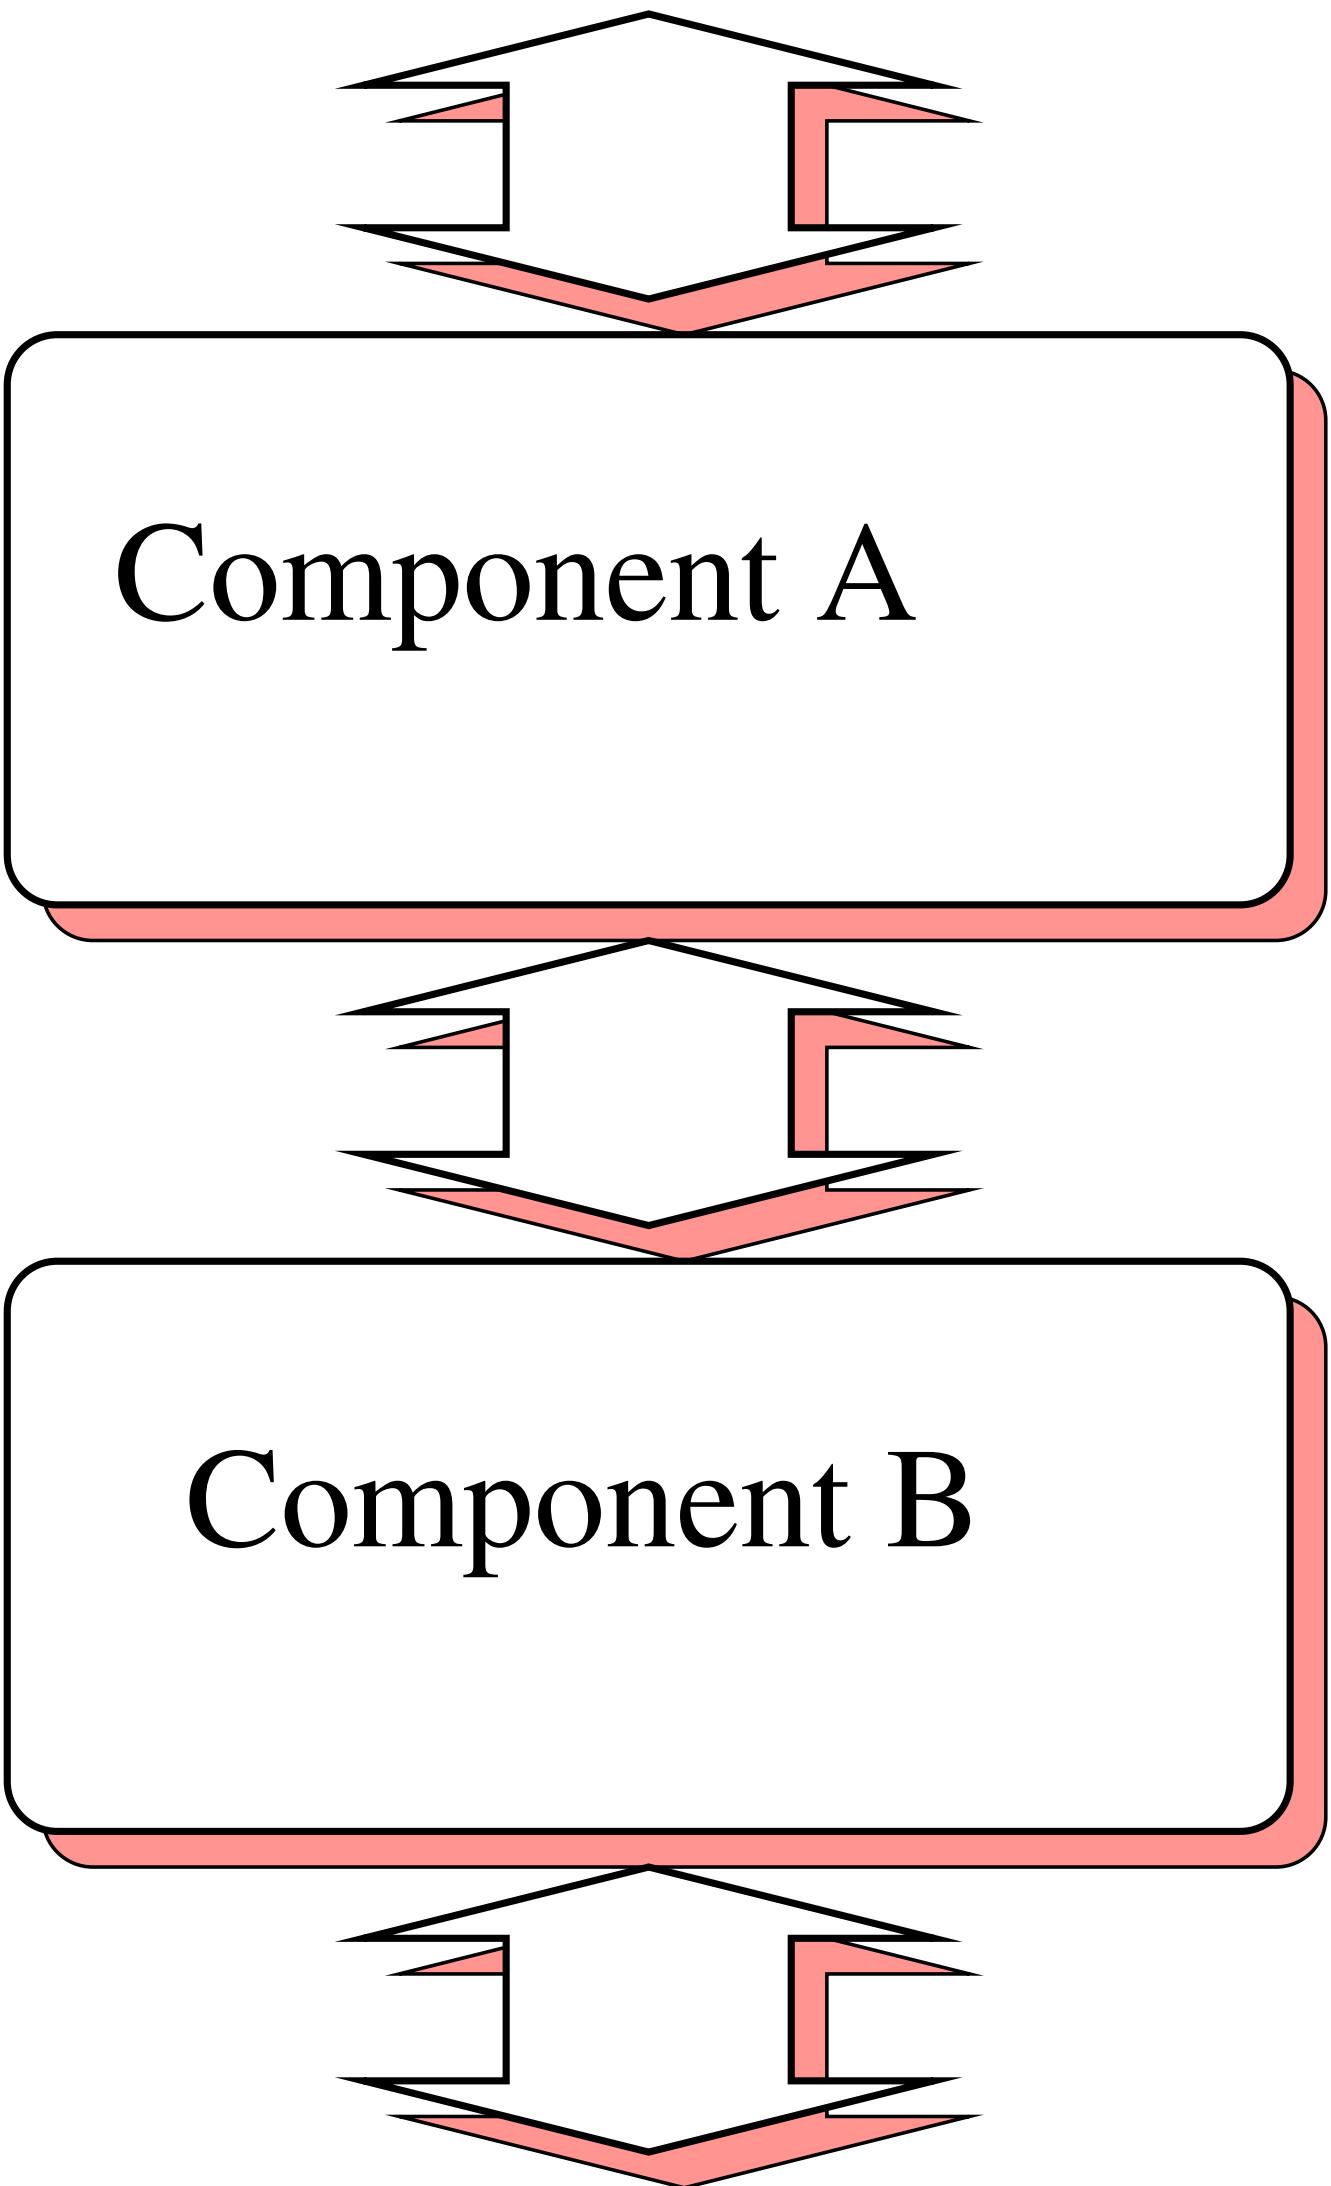
This Course

- Reliable broadcast
- Causal order broadcast
- Total order broadcast
- Consensus
- Shared memory
- Atomic commit
- Leader election
- Group membership

In-process Interaction

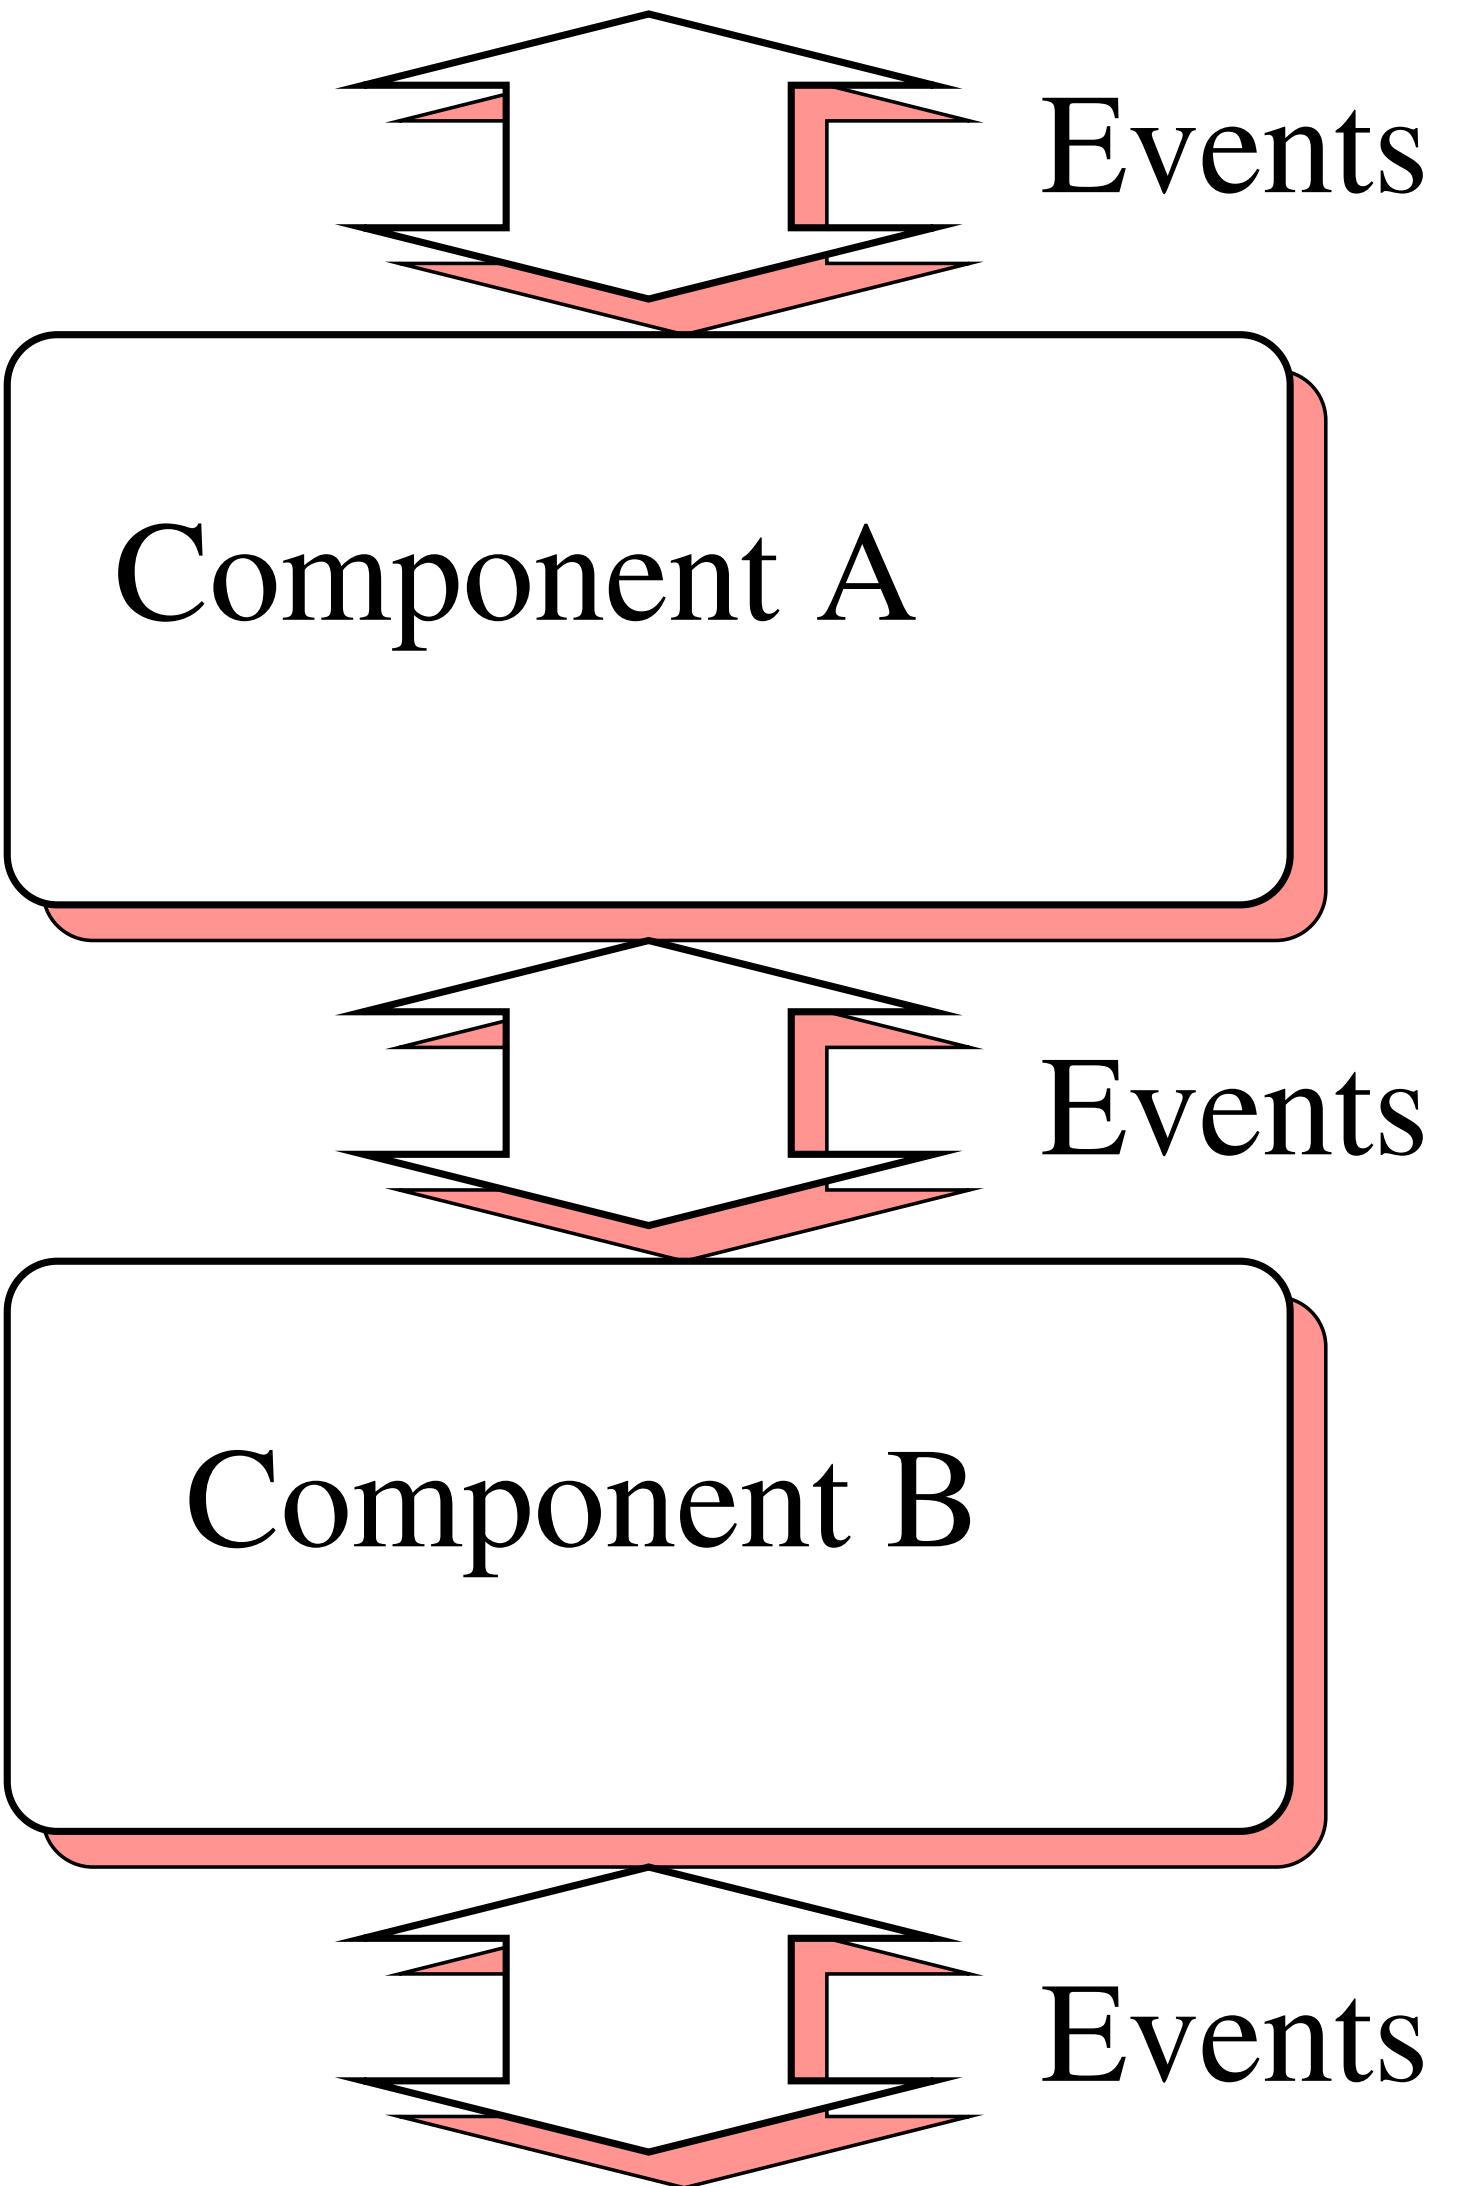
Composition Model

- The Program (or **Process**):
 - Composed of a finite set of *protocol modules*
- Organized into a stack or graph
 - Layer = Module = Component



Reactive Computing Model

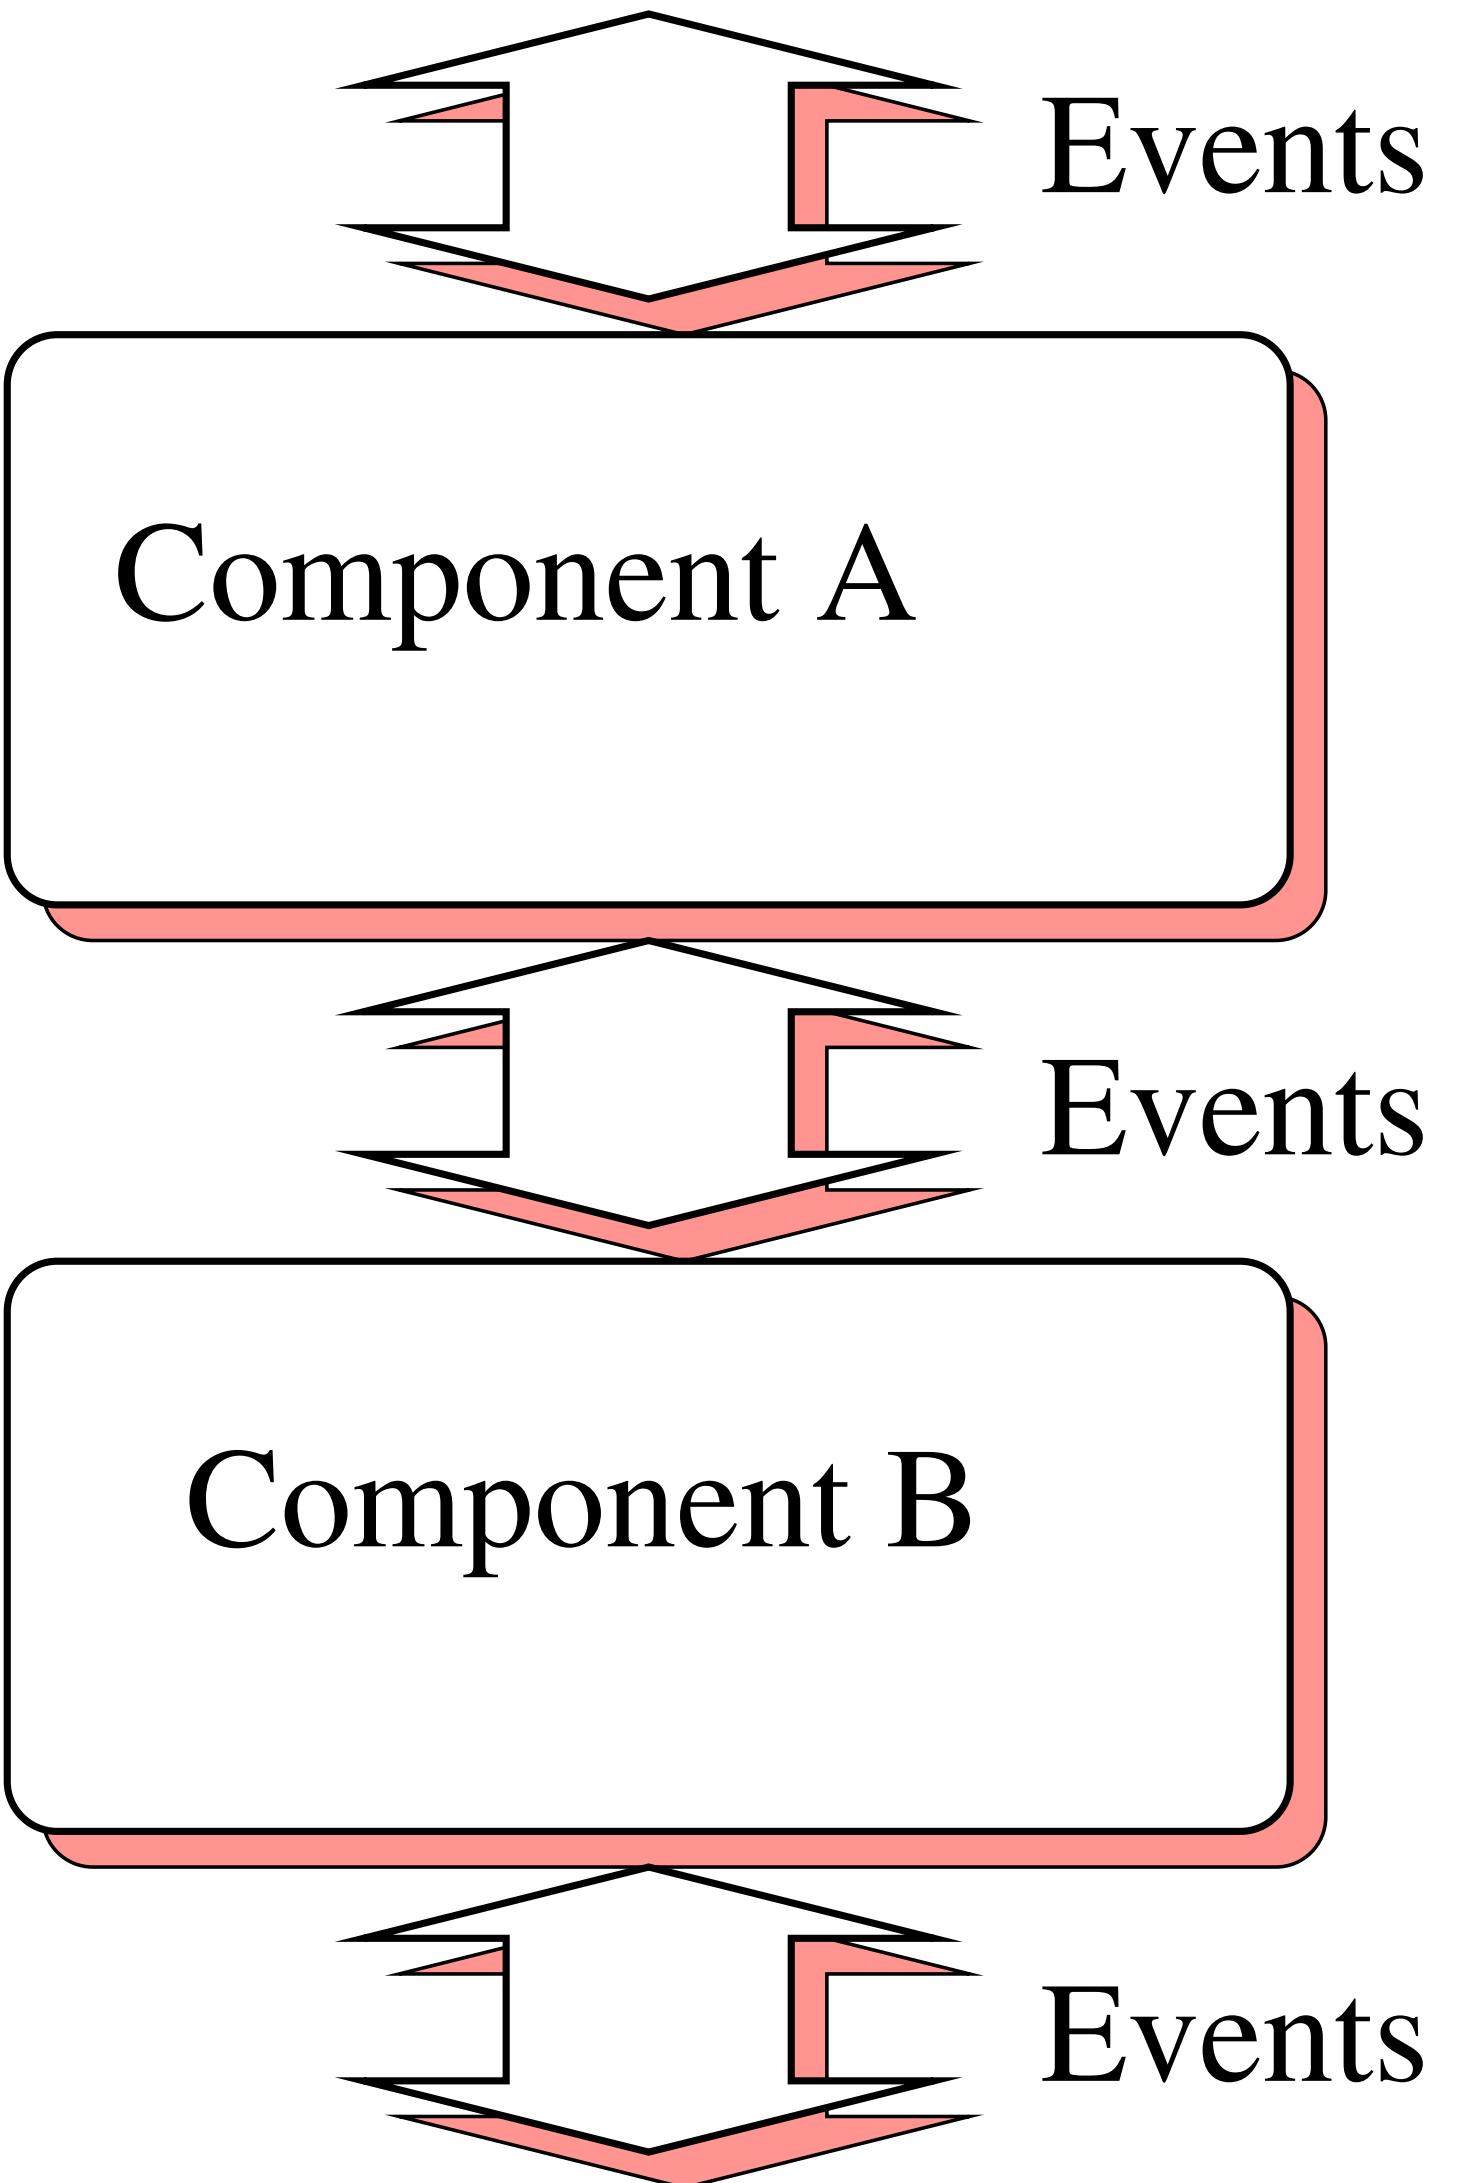
- Modules (*same process*)
 - Interact by **exchanging events**
- Algorithm (implemented by a *module*)
 - Described as a set of **event handlers**



Reactive Computing Model

- Modules (*same process*)
 - Interact by **exchanging events**
- Algorithm (implemented by a *module*)
 - Described as a set of **event handlers**

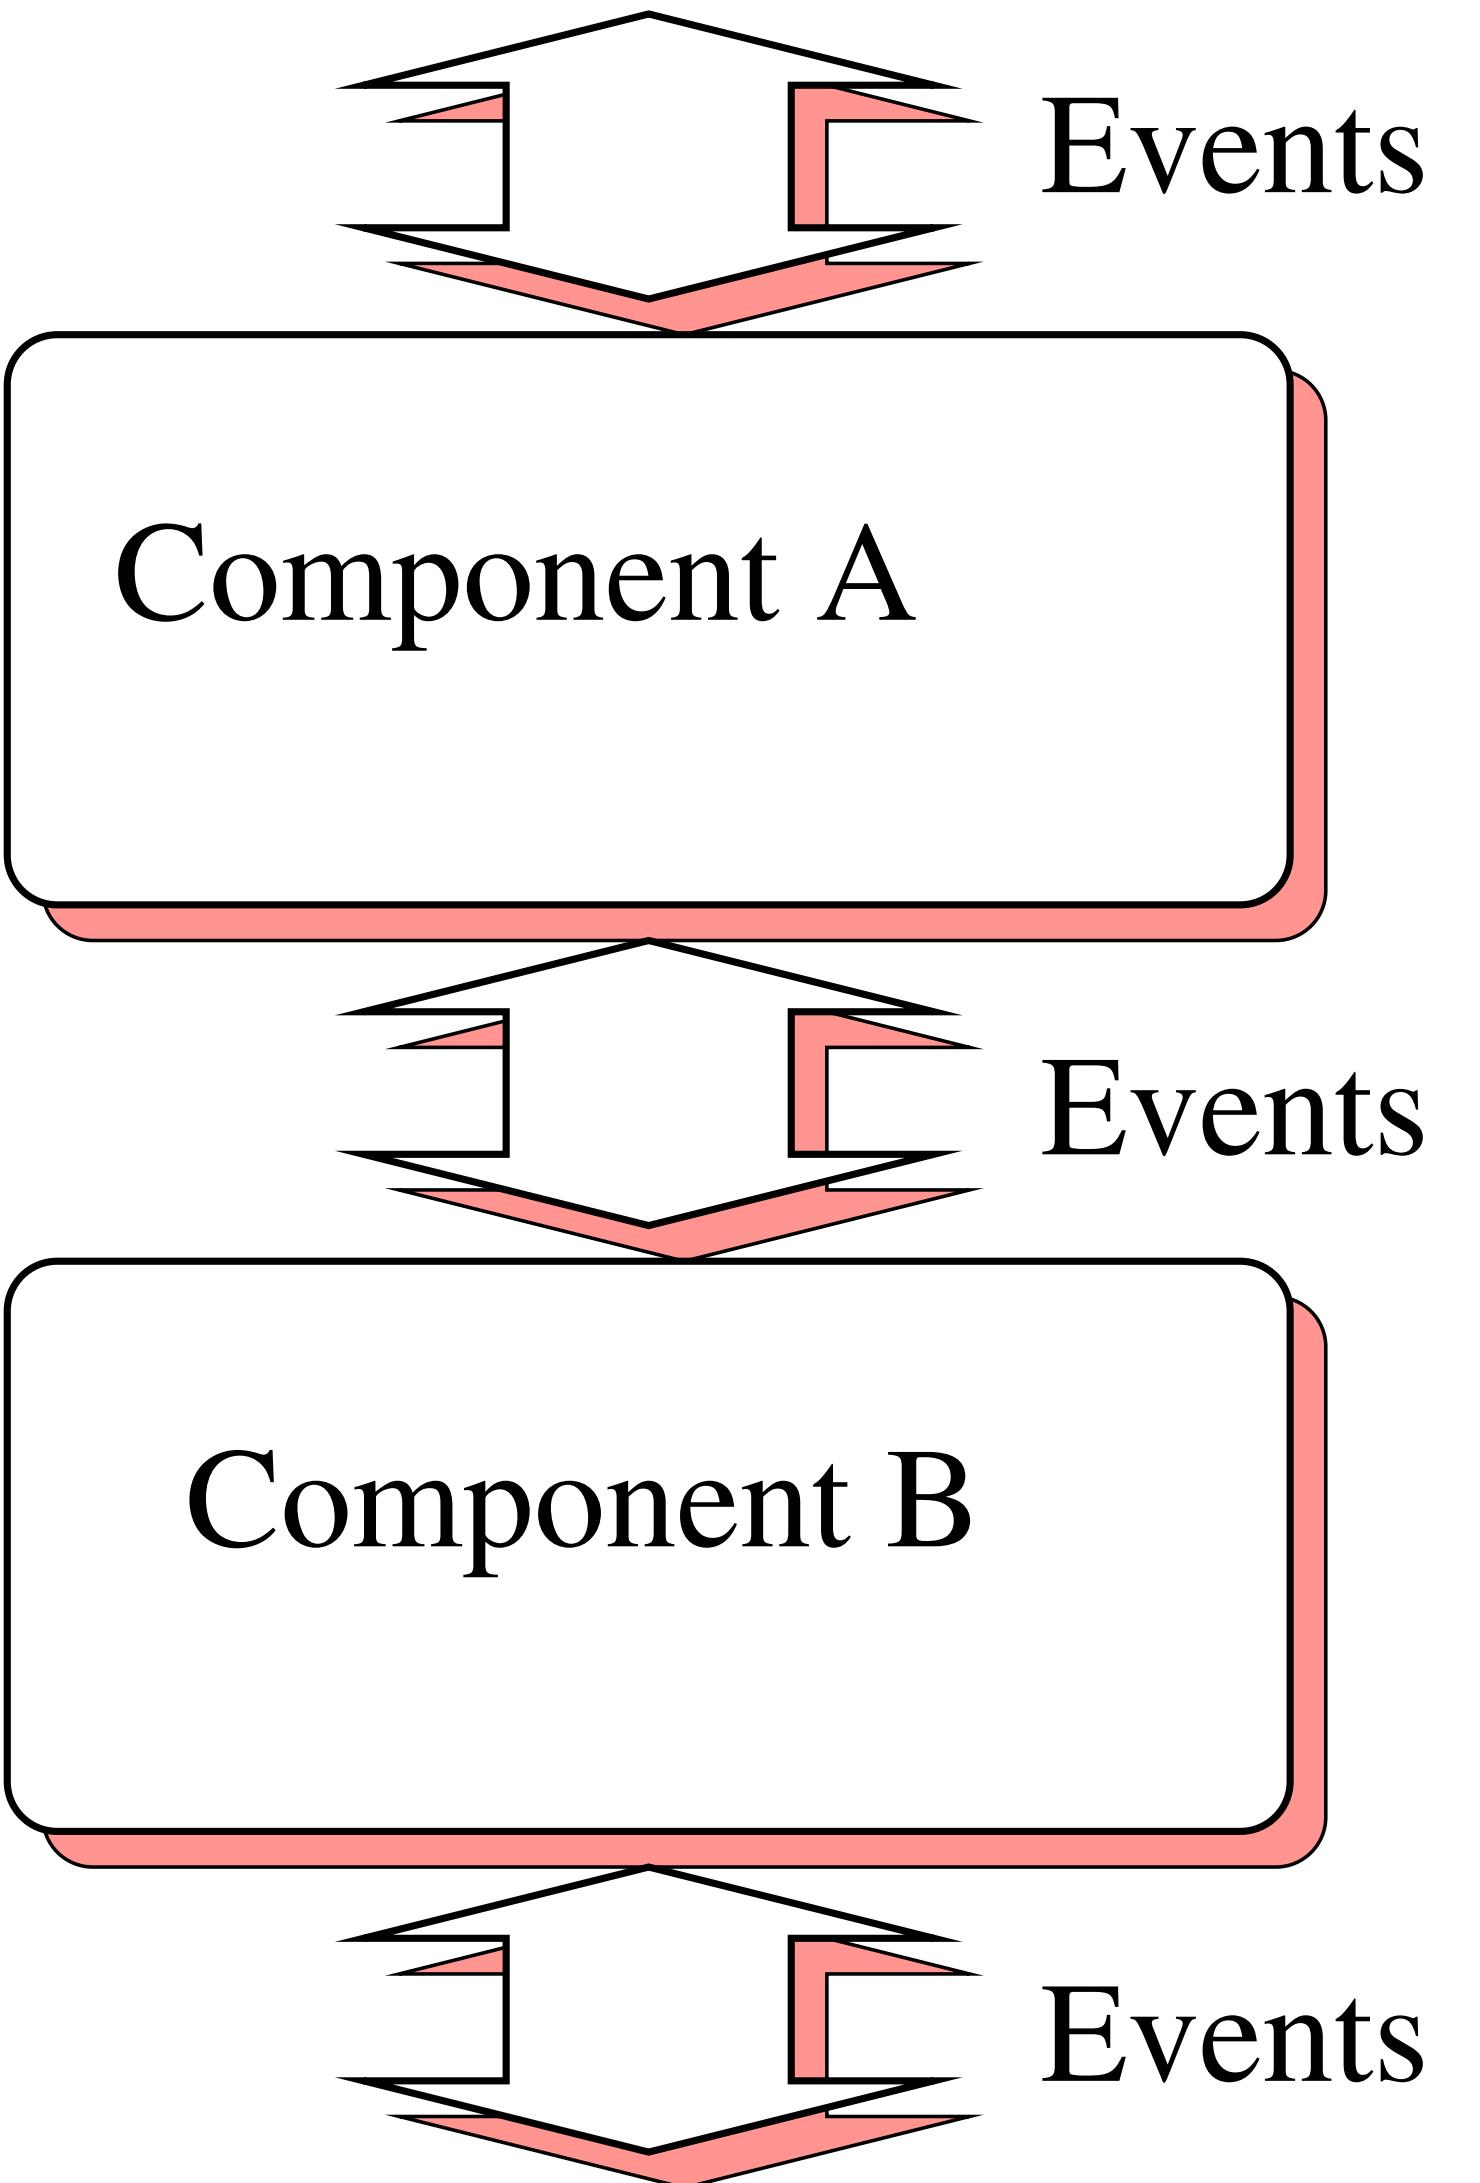
upon event $\langle co_1, Event_1 \mid att_1^1, att_1^2, \dots \rangle$ **do**
 do something;
 trigger $\langle co_2, Event_2 \mid att_2^1, att_2^2, \dots \rangle$;



Event Handlers

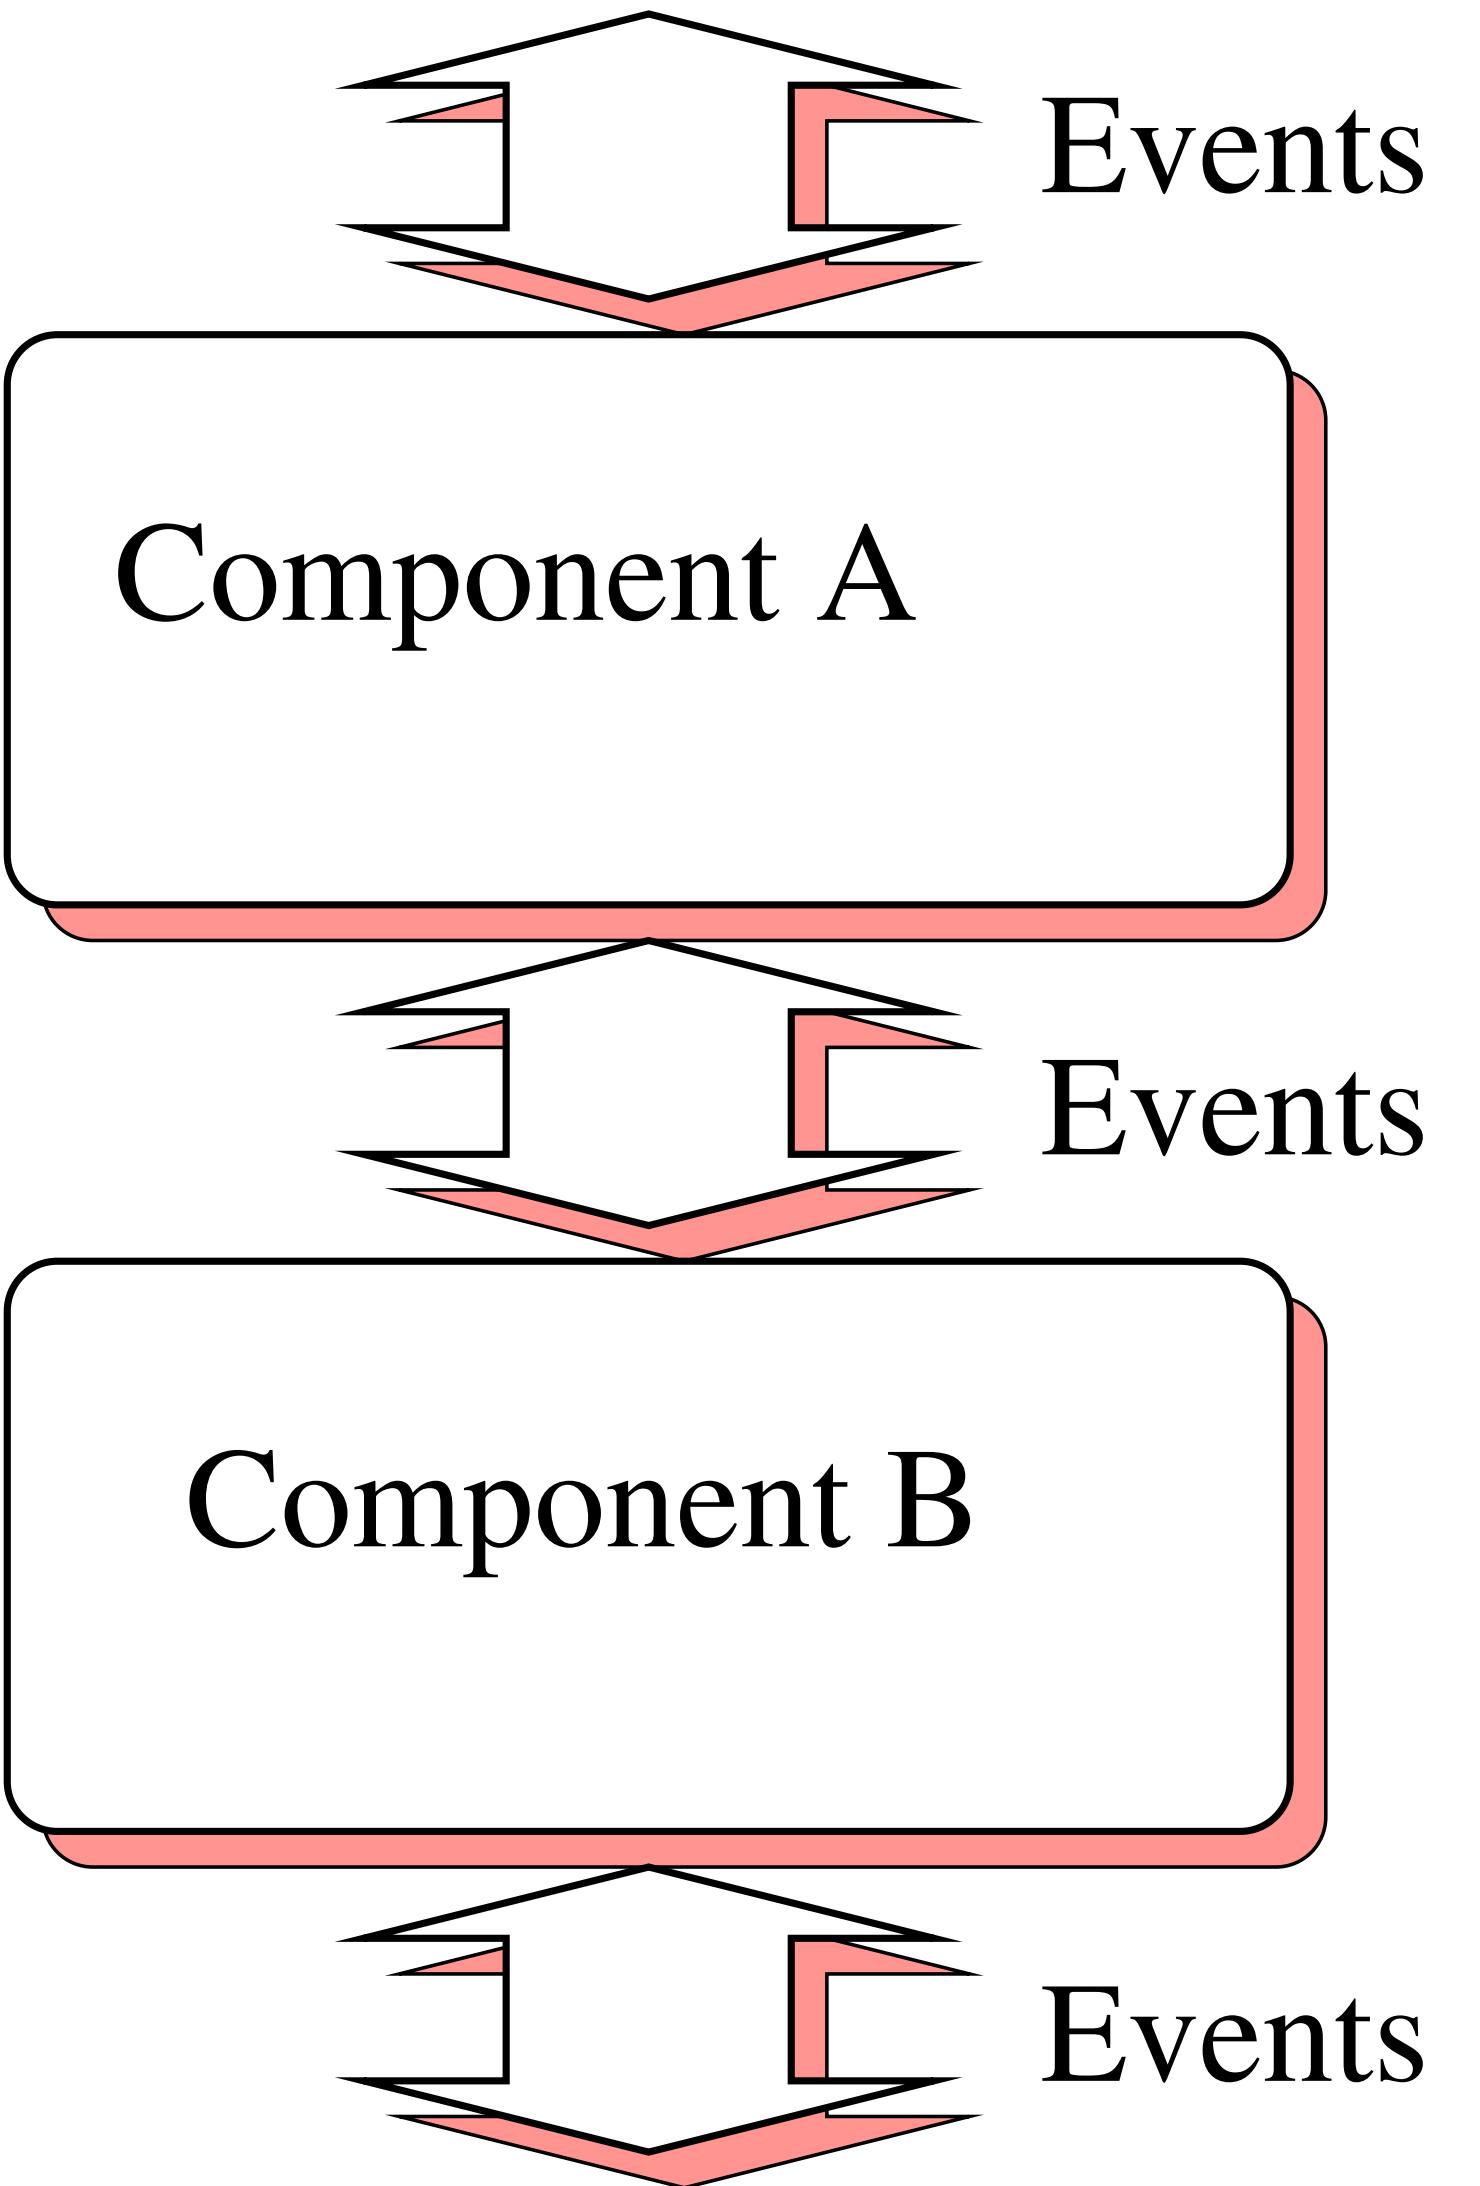
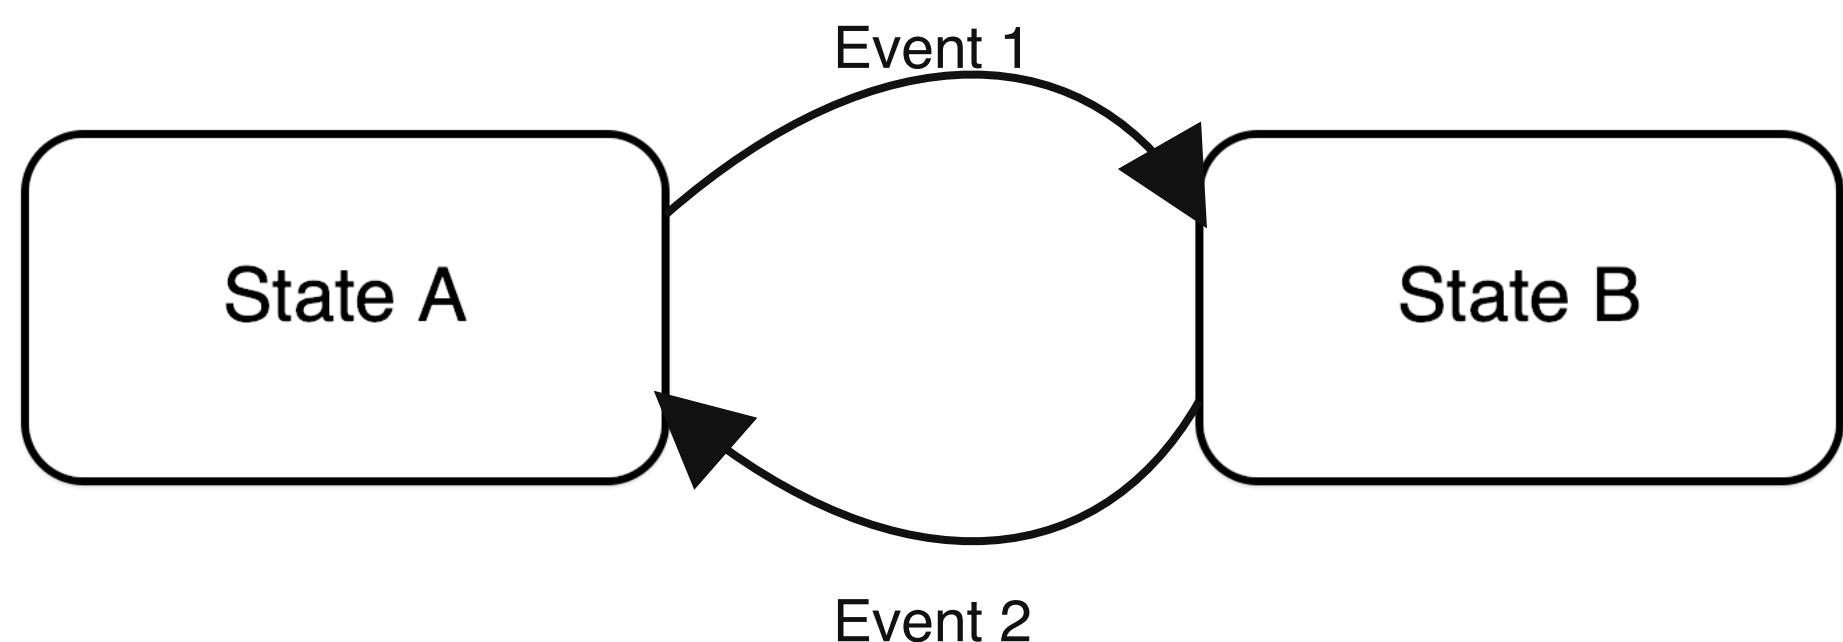
- Think of a Module as a **State Machine**
 - Events trigger state transitions
 - Incoming event / Trigger some other event

upon event $\langle co_1, Event_1 \mid att_1^1, att_1^2, \dots \rangle$ **do**
 do something;
 trigger $\langle co_2, Event_2 \mid att_2^1, att_2^2, \dots \rangle$;



Event Handlers

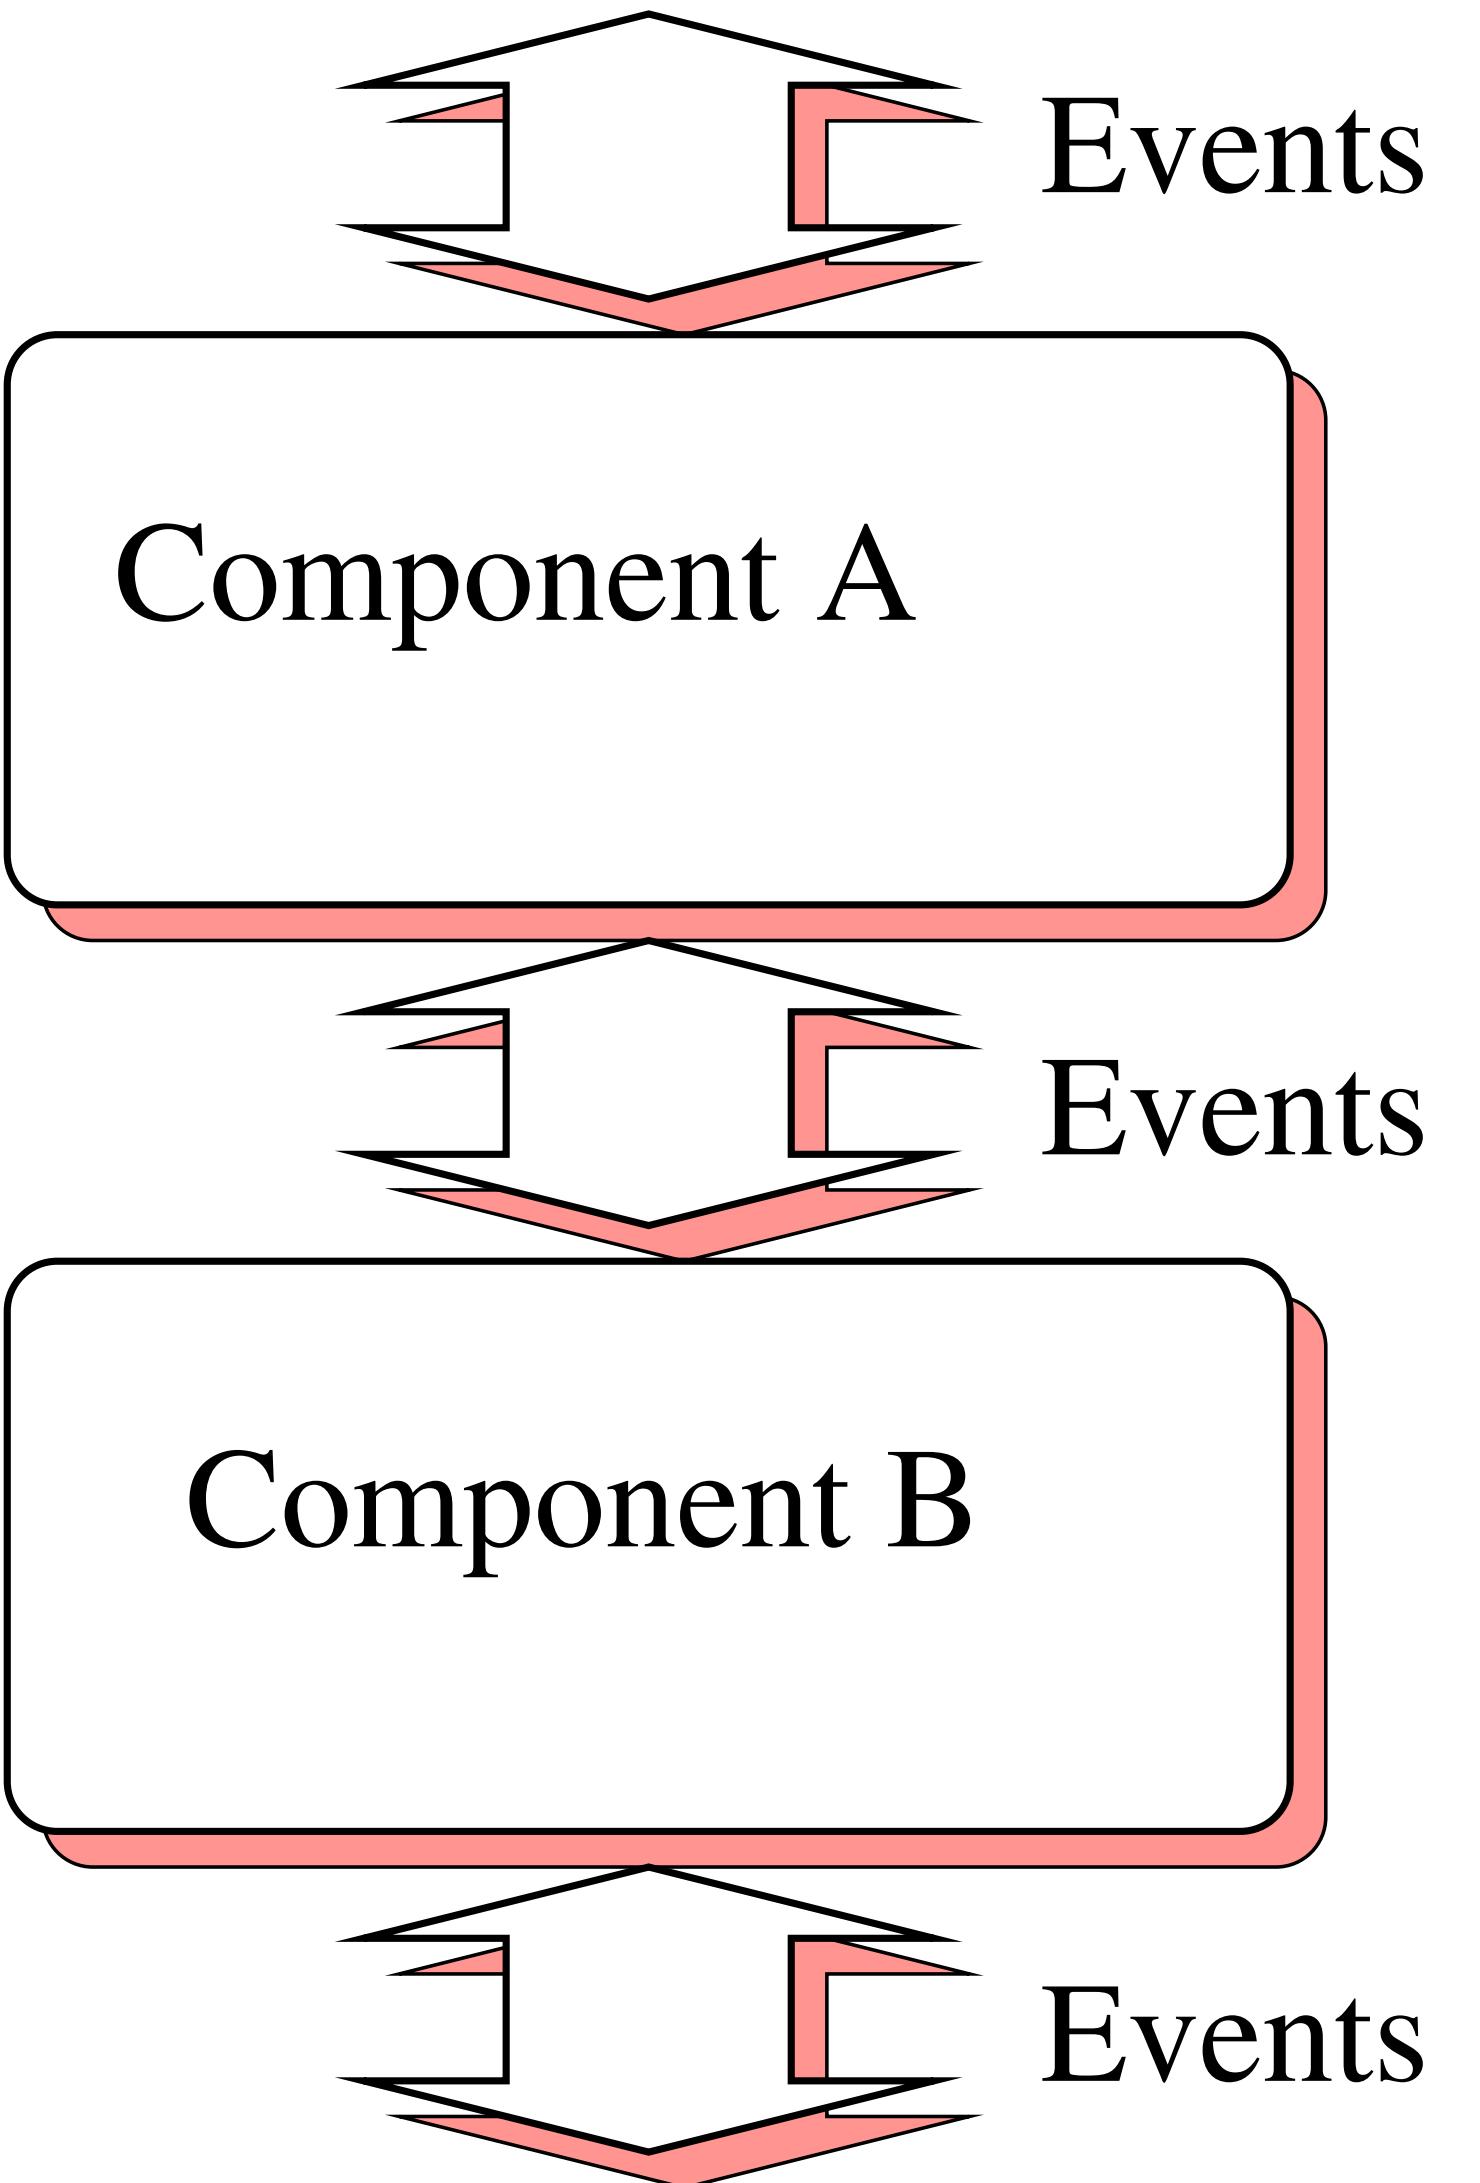
- Think of a Module as a **State Machine**
 - Events trigger *state transitions*
 - Incoming event / Trigger some other event



Event Handlers

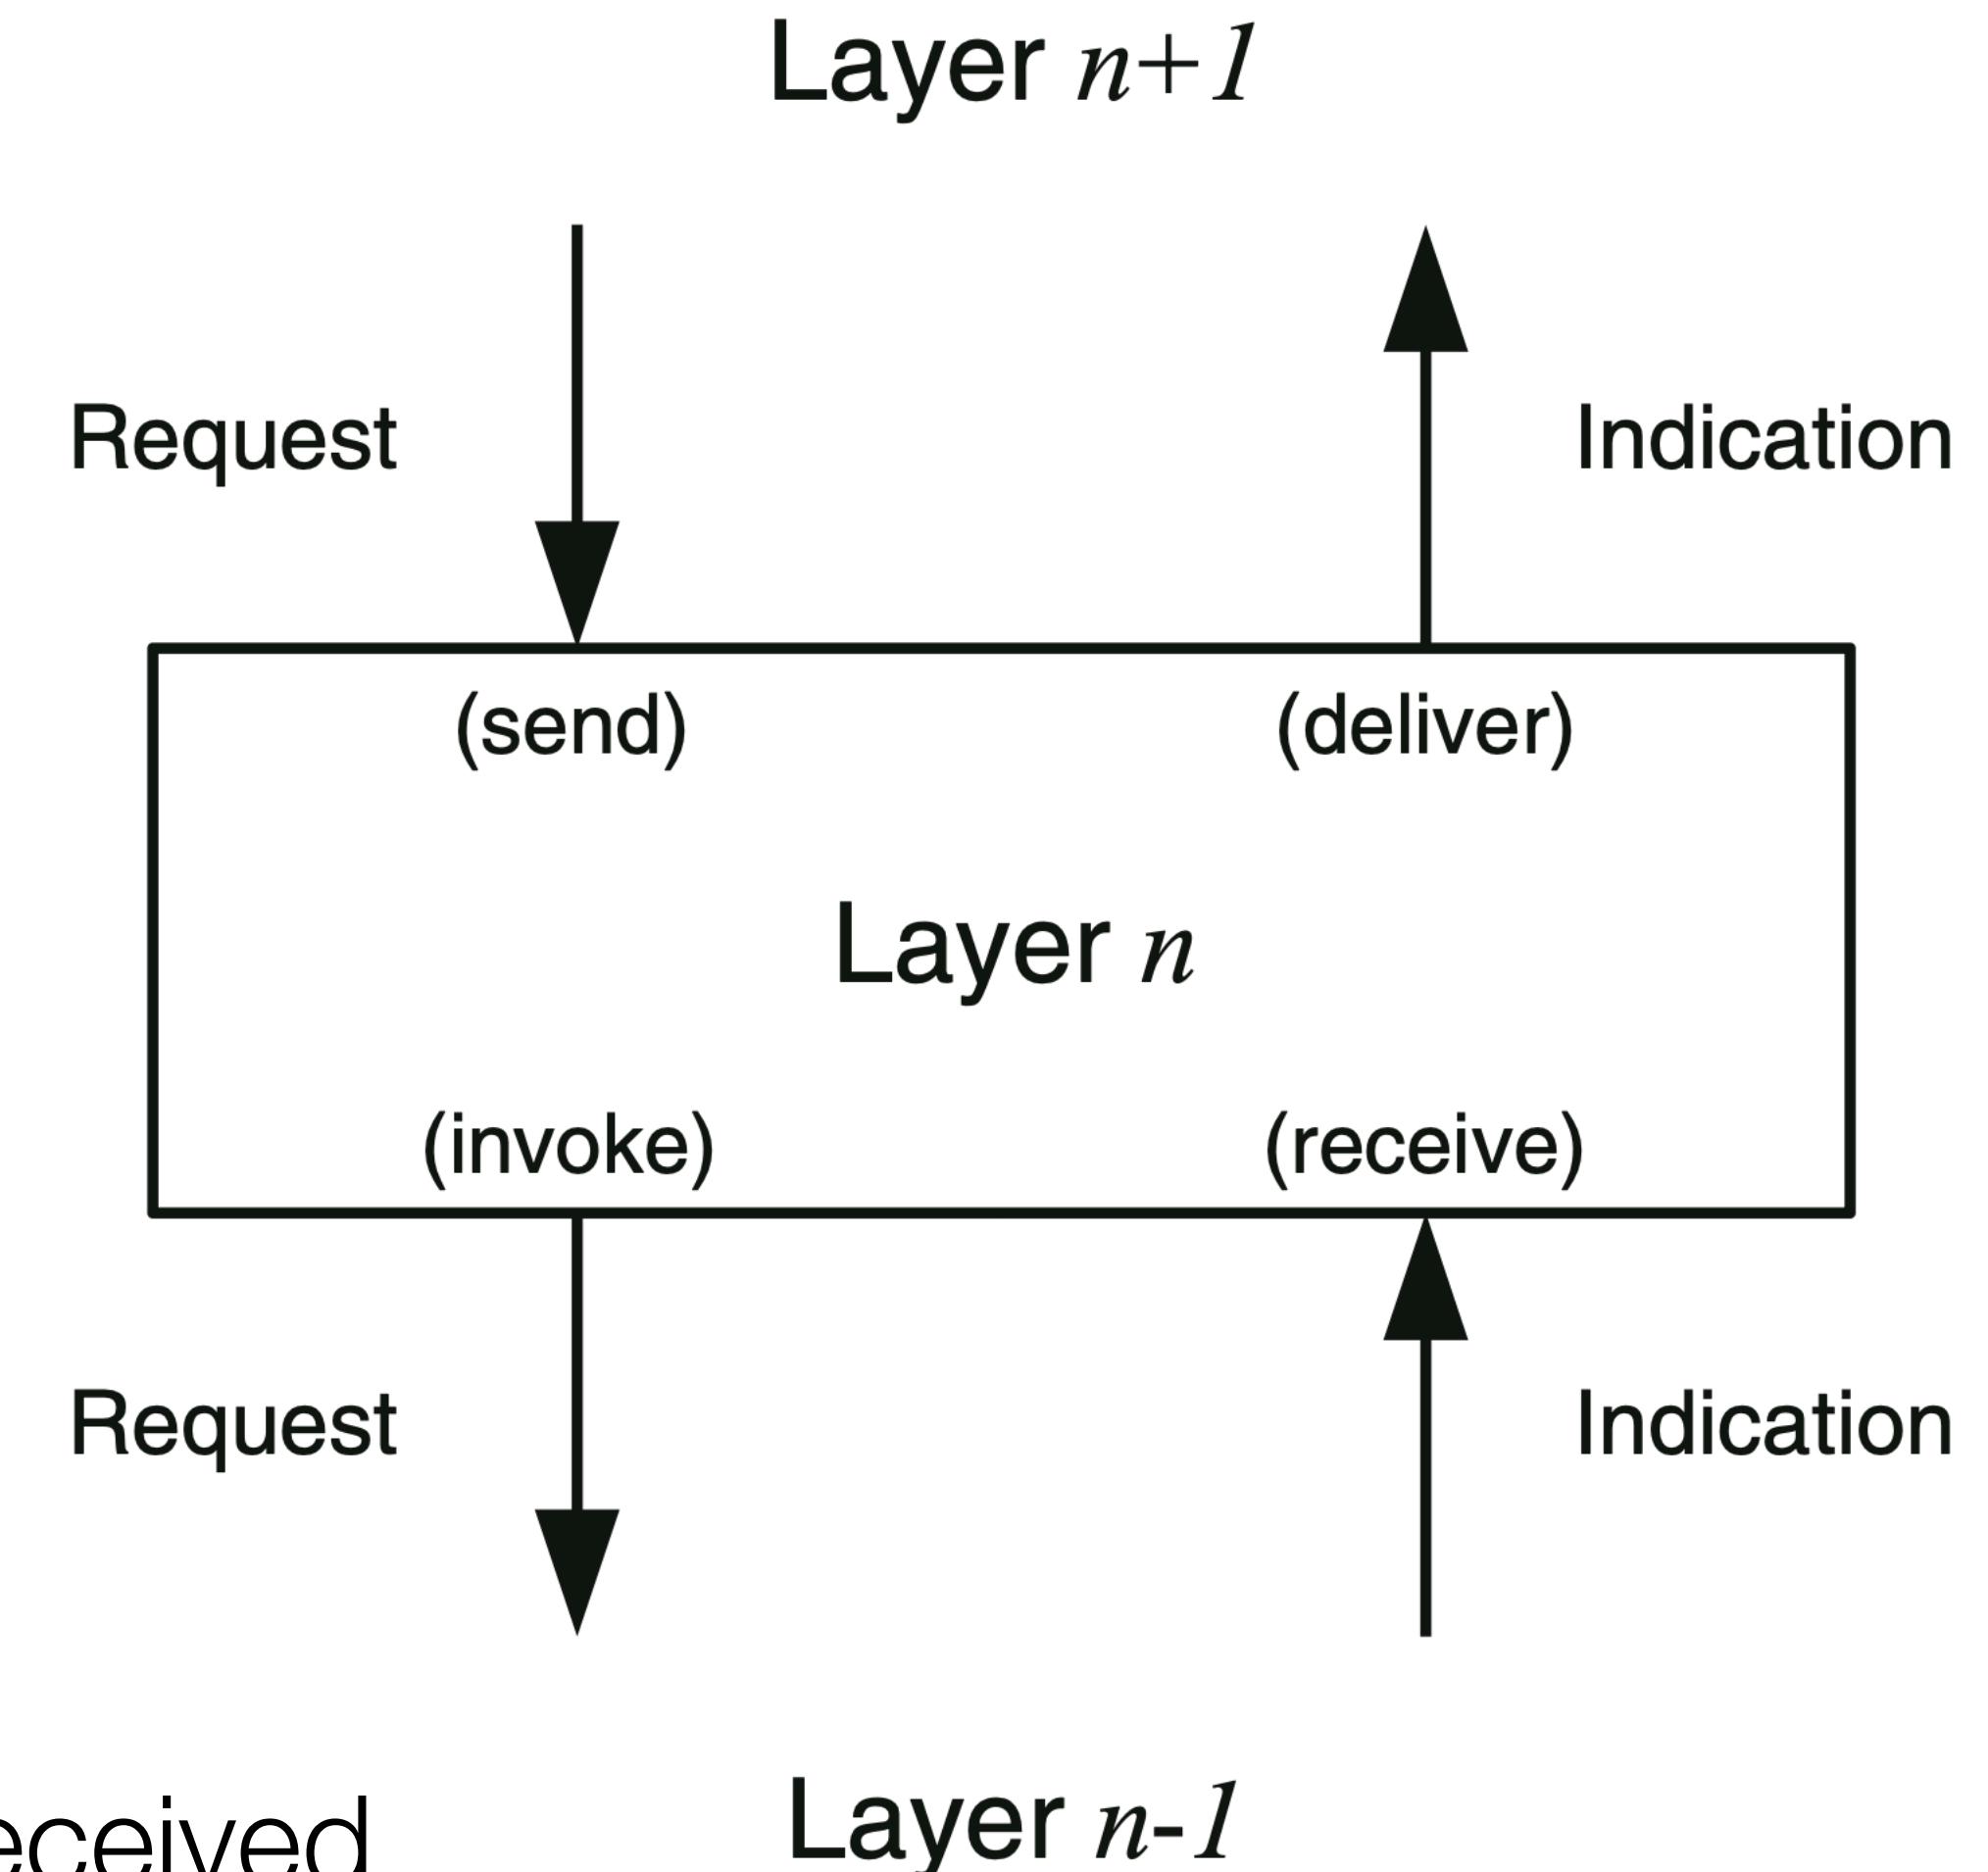
- Events may carry information in *attributes*
- Events are processed *atomically*, i.e.,
 - no two events are processed concurrently

upon event $\langle co_1, Event_1 \mid att_1^1, att_1^2, \dots \rangle$ **do**
 do something;
 trigger $\langle co_2, Event_2 \mid att_2^1, att_2^2, \dots \rangle$;



Programming Interface

- **Request** events (downcall)
 - ask for service from another module
 - e.g., to broadcast a message
- **Indication** events (upcall)
 - used to deliver information to another module
 - e.g., to deliver a message in-order that was received out-of order

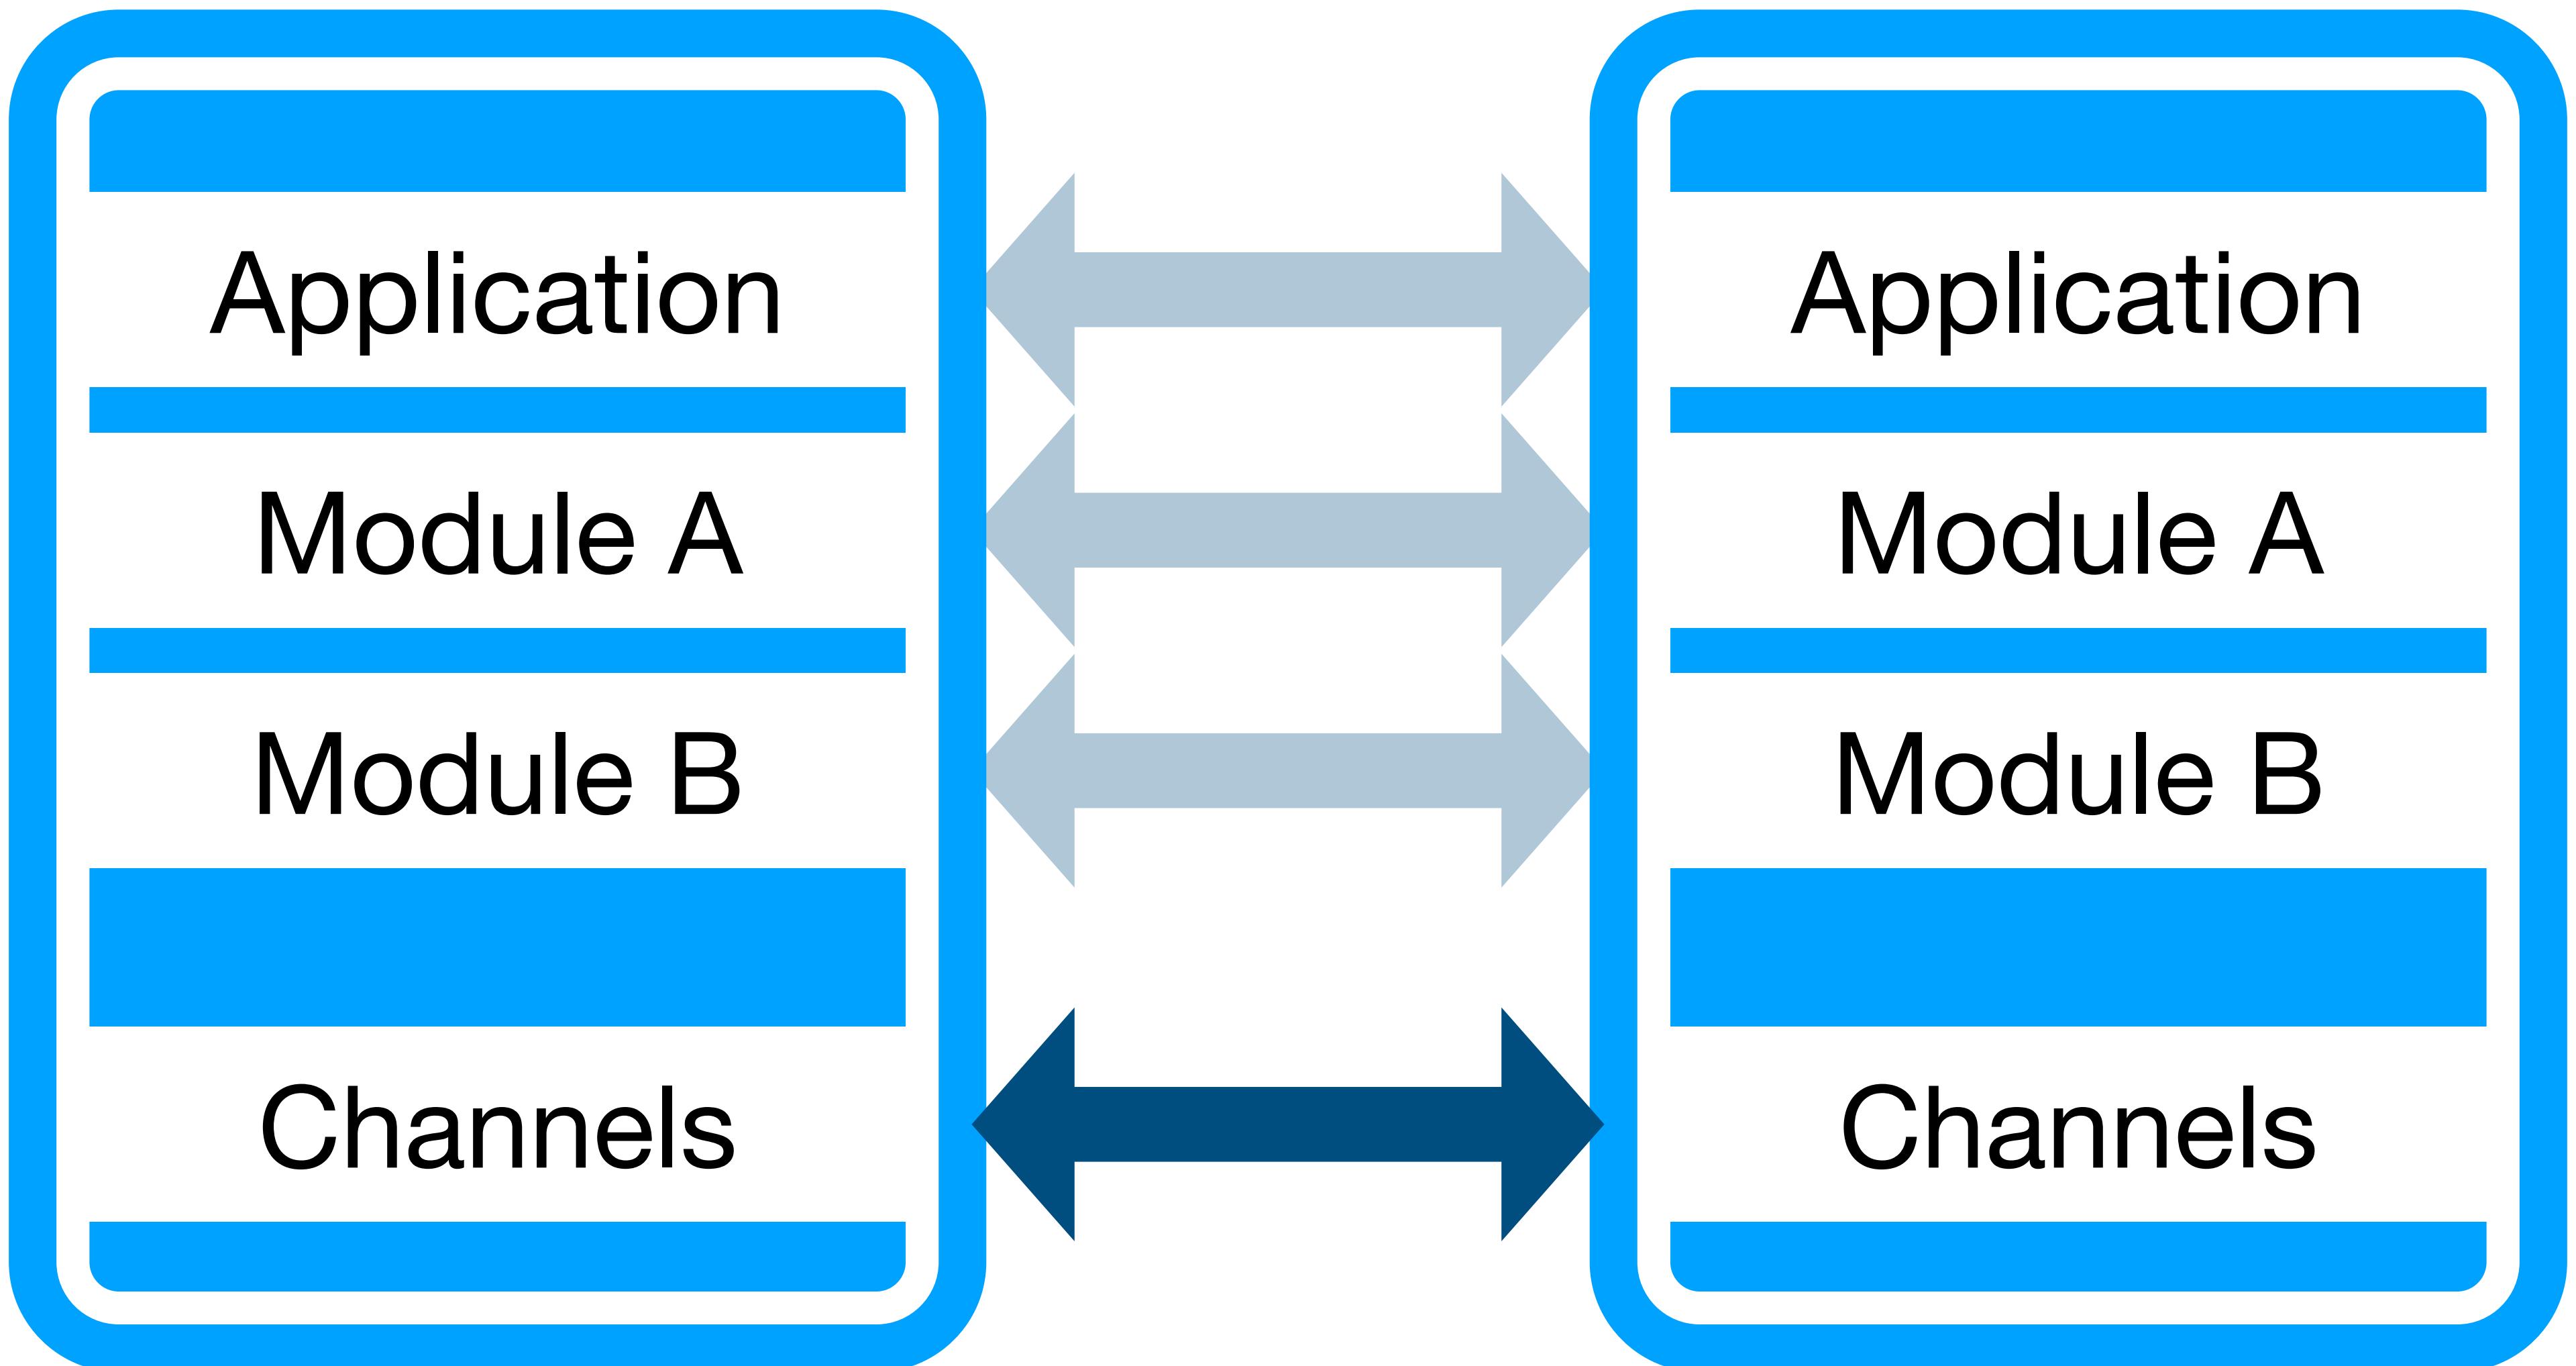


Interaction Between Processes

Module Interaction Between Processes

- Modules interact with corresponding modules on **peer processes**
- May have multiple instances (copies) of a module in one process

1.4 Software Components



Communication Between Modules

- Asynchronous interaction
 - **non-blocking**: send request *don't wait* for reply/result
- Synchronous interaction
 - **blocking**: send request *wait* for reply/result

Textbook

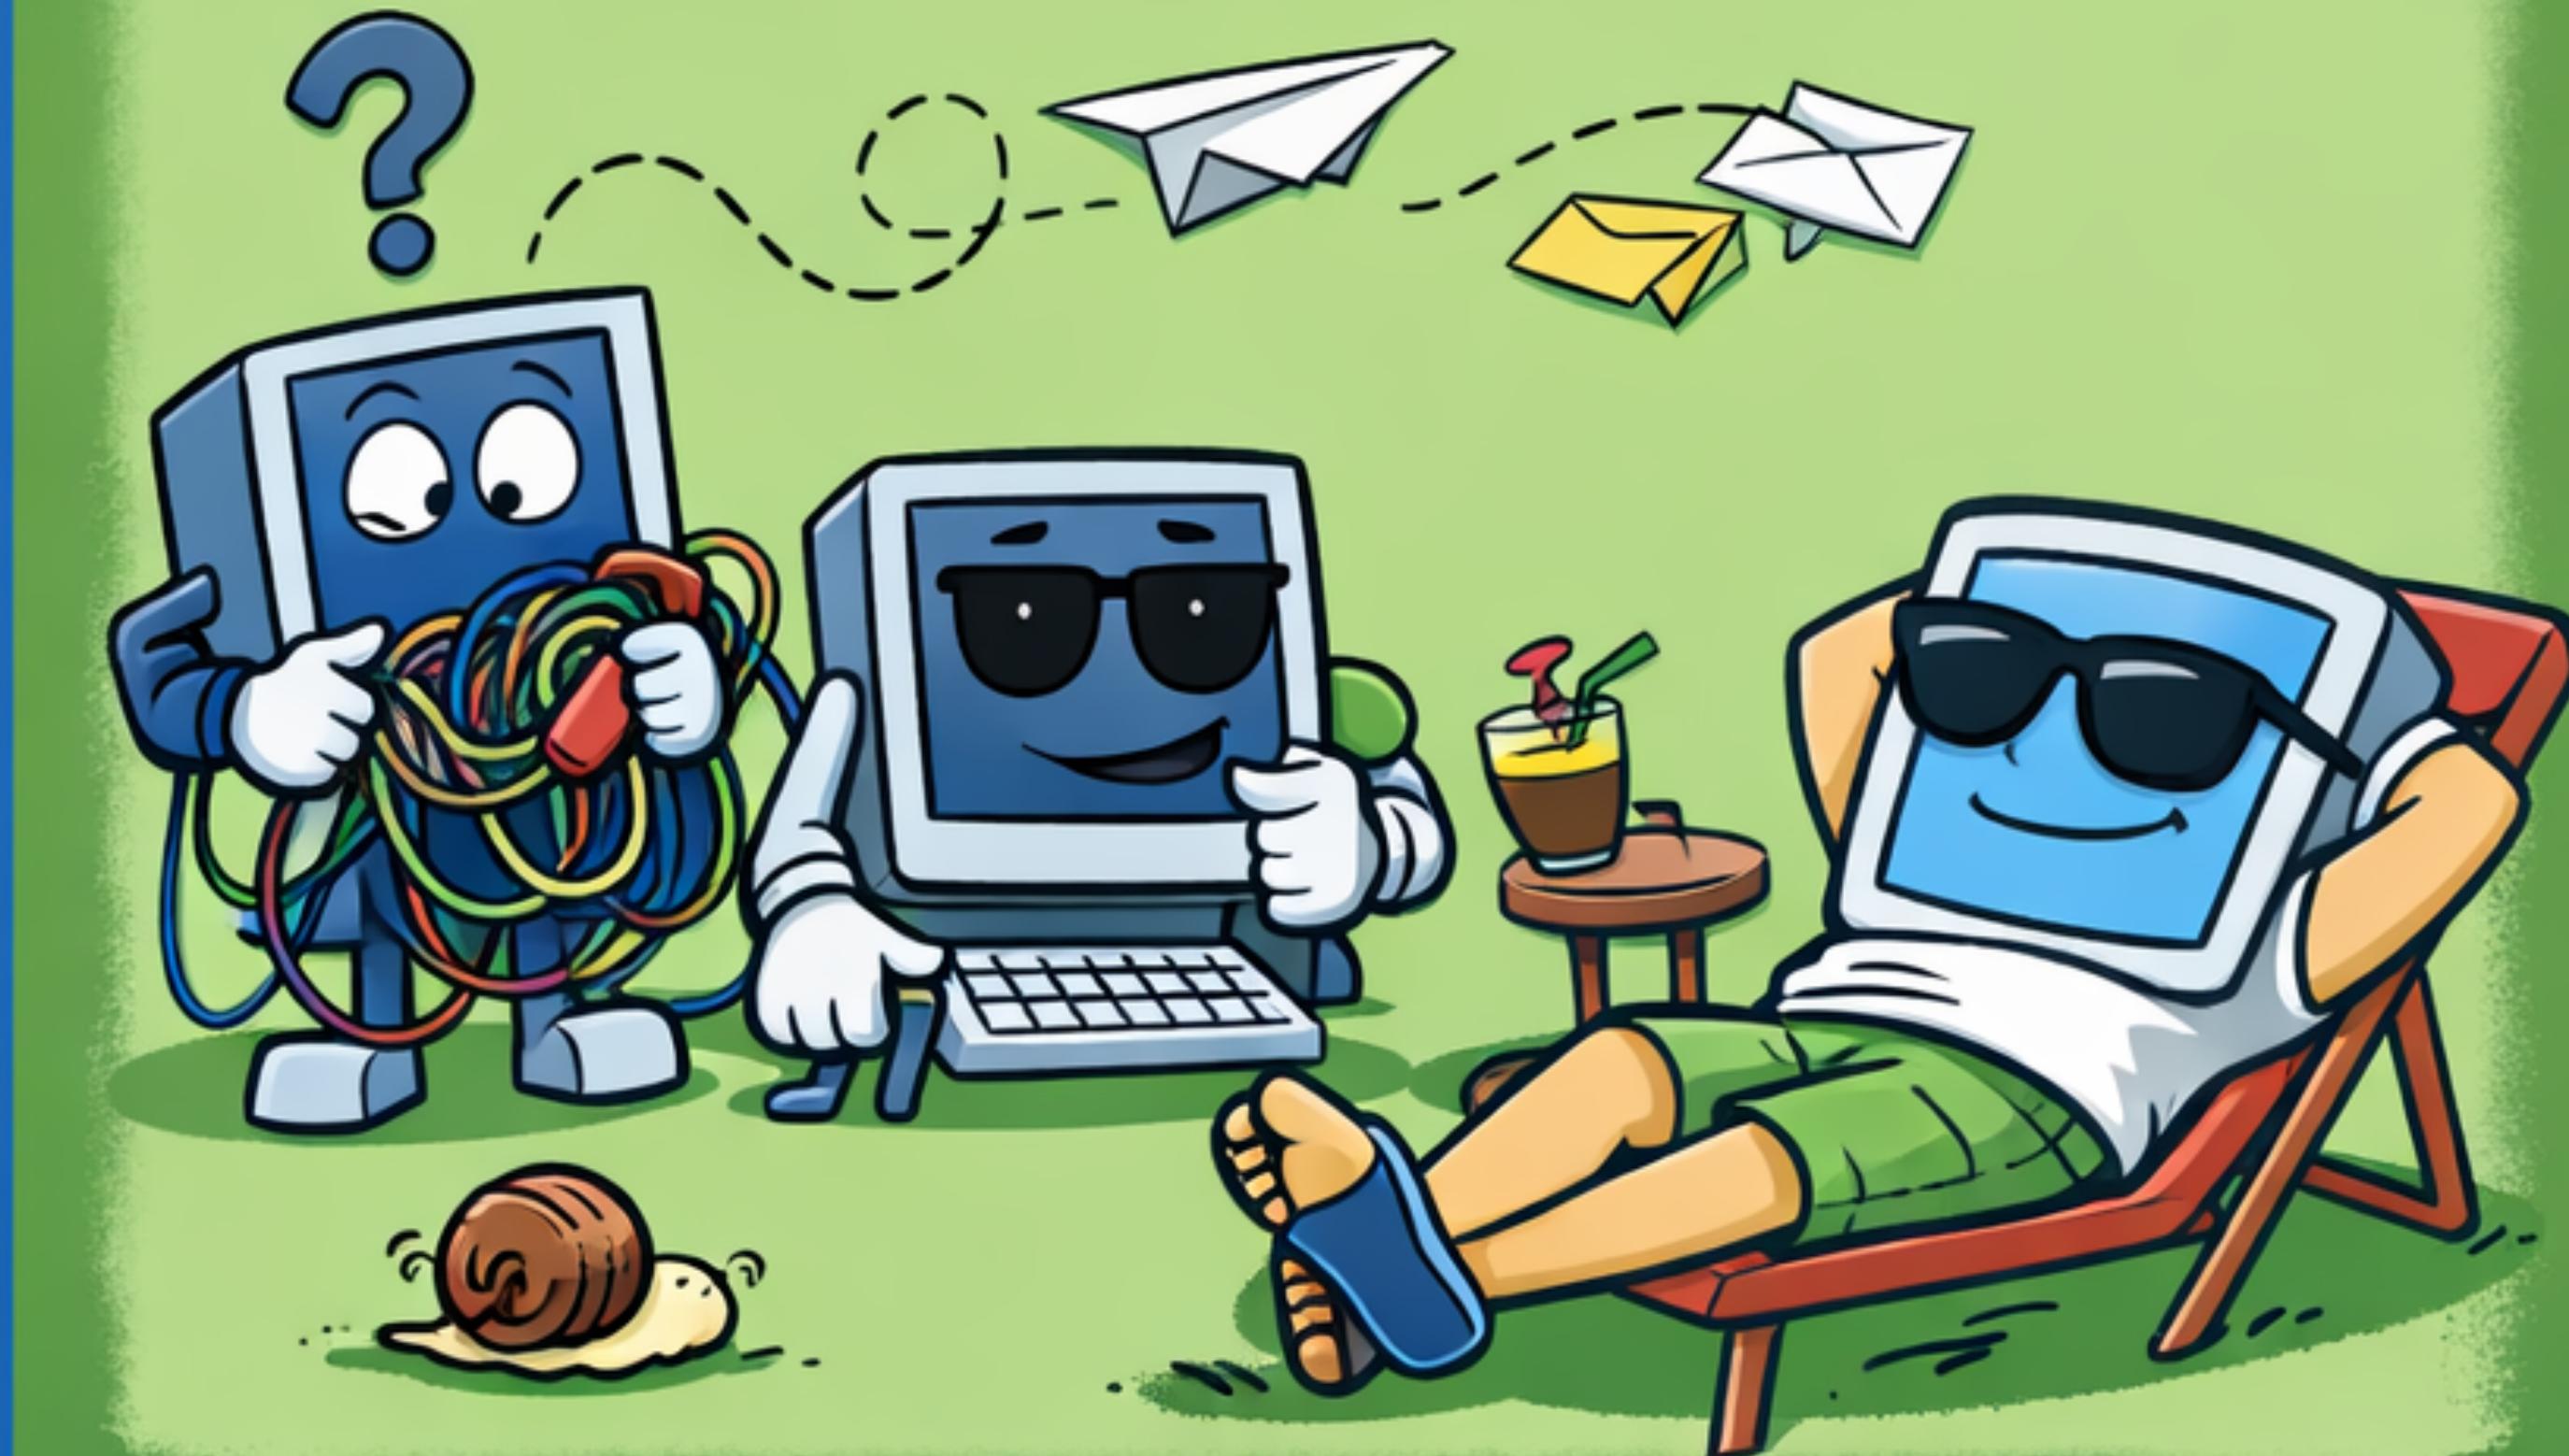
Joke time!

Synchronous



We're all just waiting for Bob..."

Asynchronous



Bob will get back to us... maybe."

Algorithmic Terminology

Protocol Module

Examples

Module 1.1: Interface and properties of a job handler

Module:

Name: JobHandler, instance jh .

Events:

Request: $\langle jh, Submit \mid job \rangle$: Requests a job to be processed.

Indication: $\langle jh, Confirm \mid job \rangle$: Confirms that the given job has been (or will be) processed.

Properties:

JH1: *Guaranteed response:* Every submitted job is eventually confirmed.

Algorithm 1.1: Synchronous Job Handler

Implements:JobHandler, **instance** *jh*.

upon event $\langle jh, Submit \mid job \rangle$ **do**
 process(*job*);
 trigger $\langle jh, Confirm \mid job \rangle$;

Algorithm 1.2: Asynchronous Job Handler

Implements:

JobHandler, **instance** jh .

upon event $\langle jh, \text{Init} \rangle$ **do**

$buffer := \emptyset;$

upon event $\langle jh, \text{Submit} \mid job \rangle$ **do**

$buffer := buffer \cup \{job\};$

trigger $\langle jh, \text{Confirm} \mid job \rangle;$

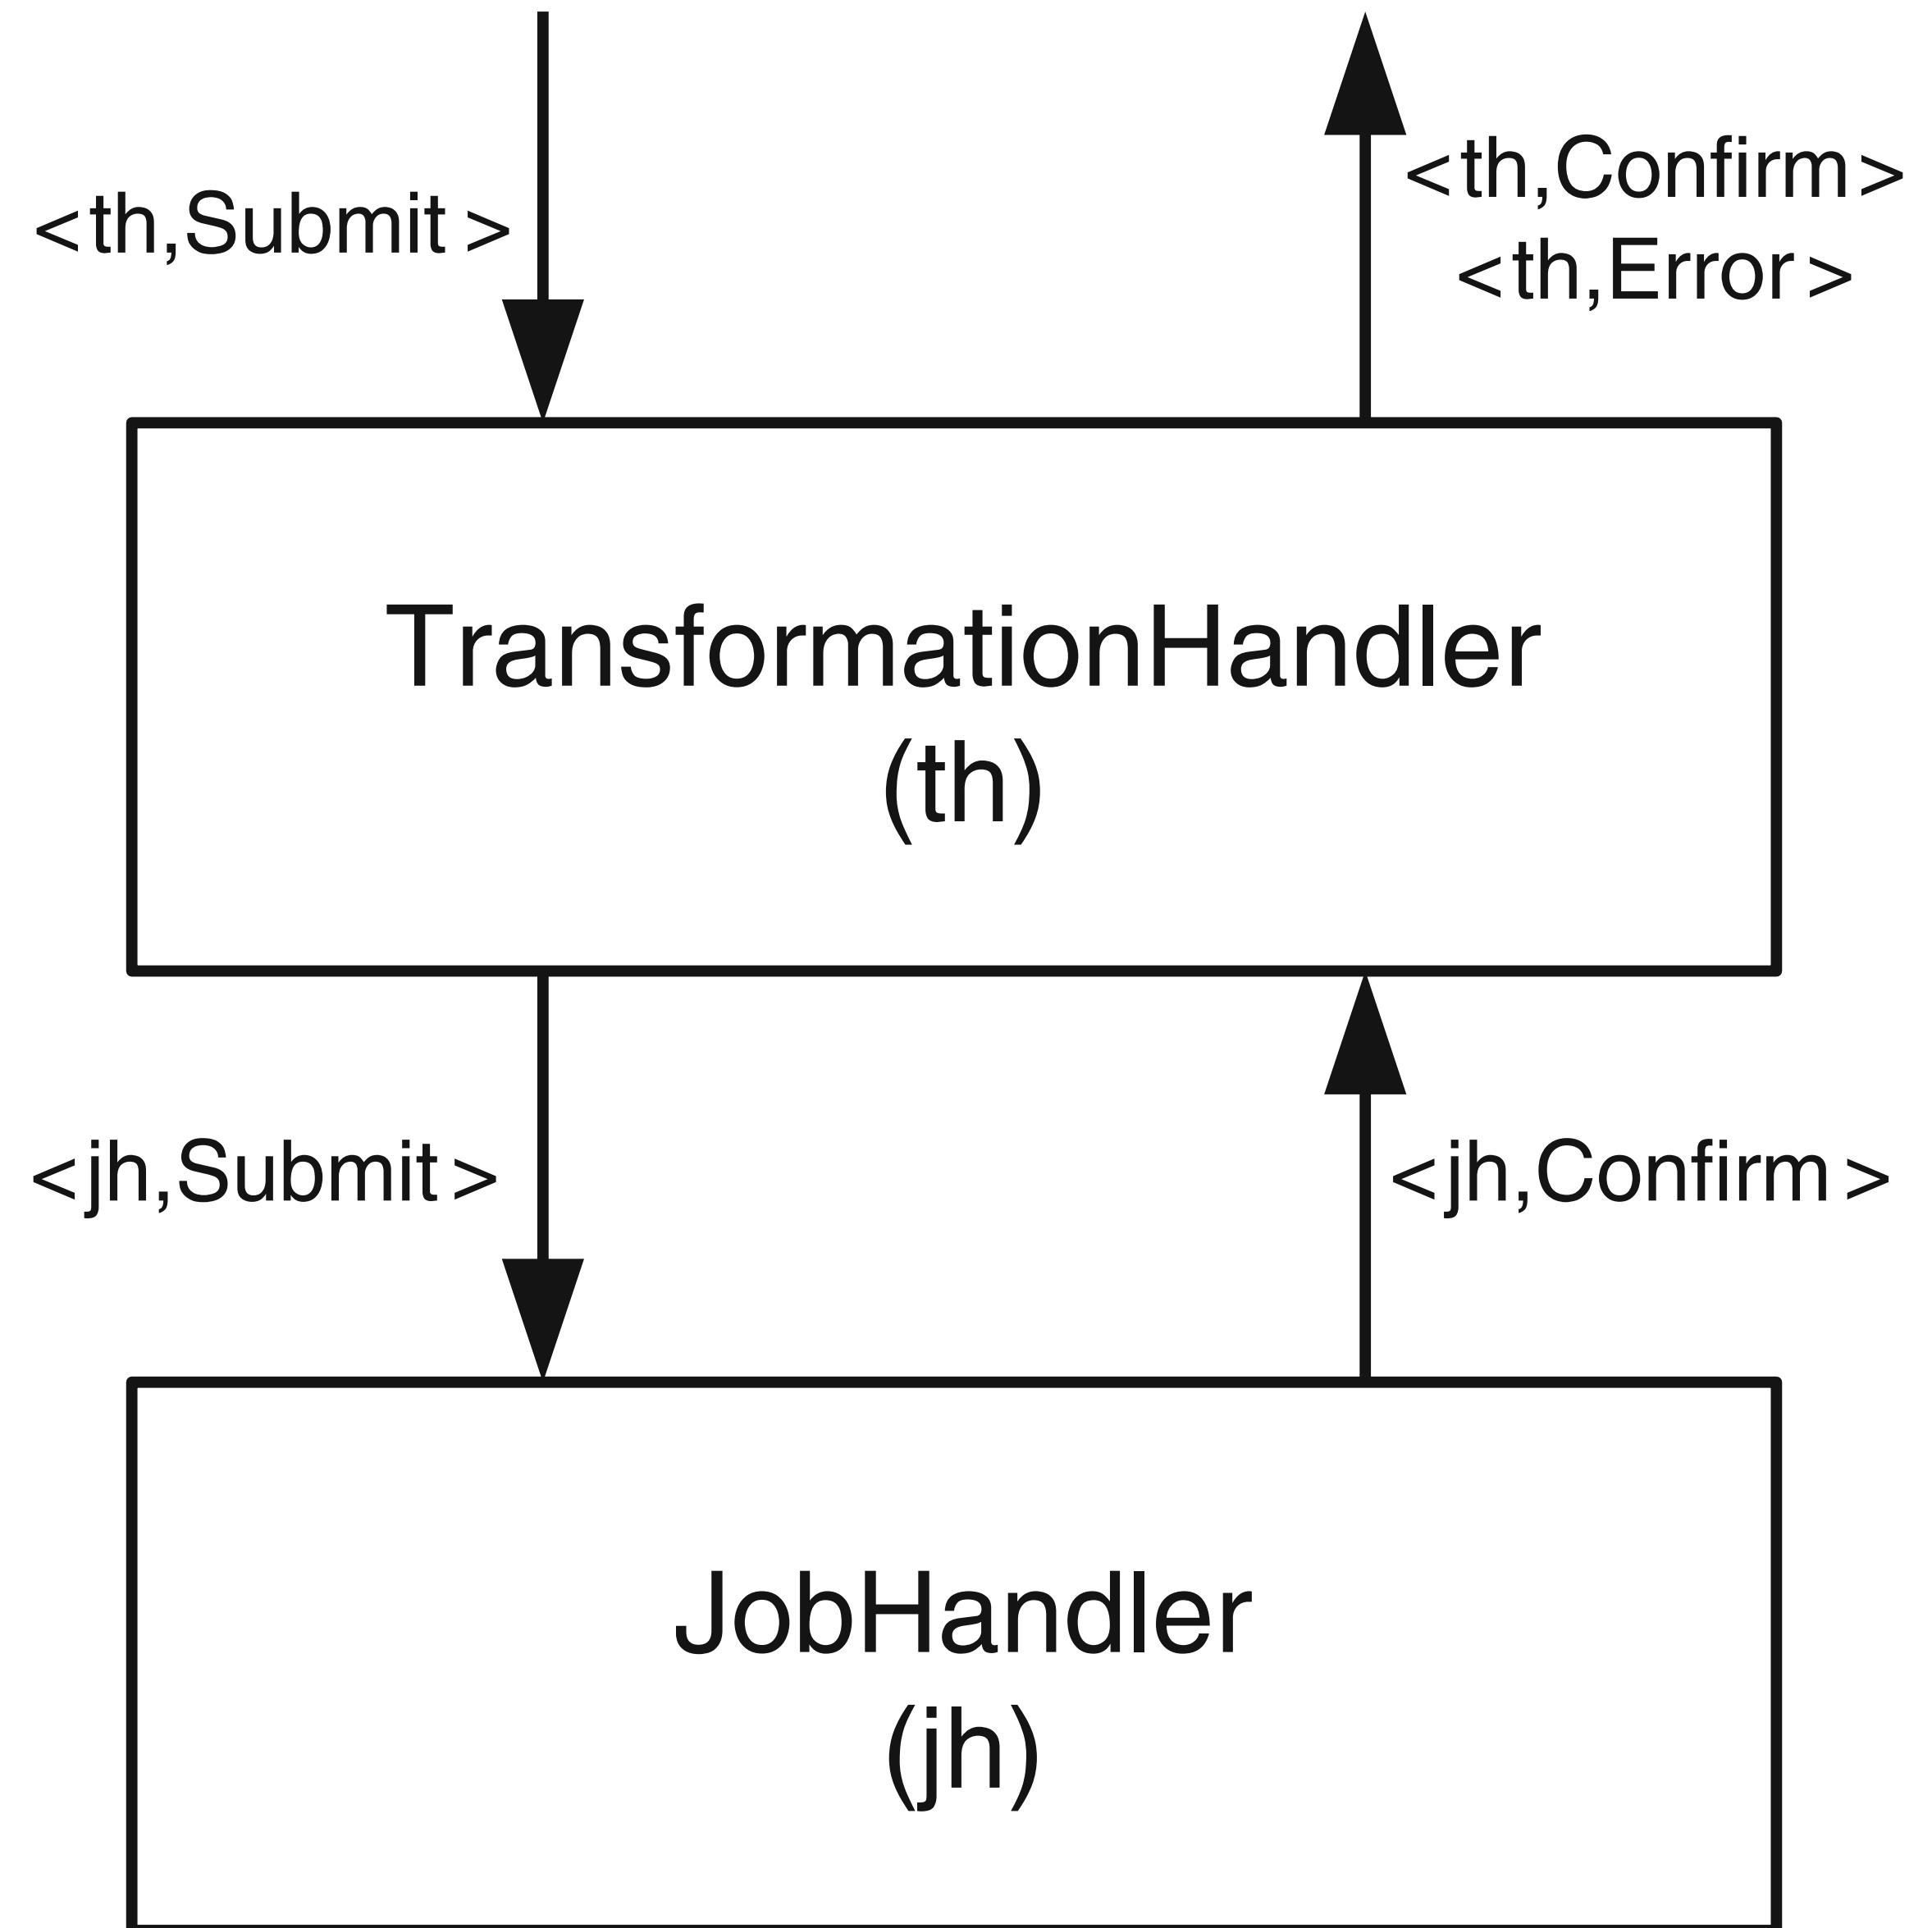
upon $buffer \neq \emptyset$ **do**

$job := \text{selectjob}(buffer);$

process(job);

$buffer := buffer \setminus \{job\};$

1.4 Software Components



Module 1.2: Interface and properties of a job transformation and processing abstraction

Module:

Name: TransformationHandler, **instance** *th*.

Events:

Request: $\langle th, Submit \mid job \rangle$: Submits a job for transformation and for processing.

Indication: $\langle th, Confirm \mid job \rangle$: Confirms that the given job has been (or will be) transformed and processed.

Indication: $\langle th, Error \mid job \rangle$: Indicates that the transformation of the given job failed.

1.4 Software Components

Algorithm 1.3: Job-Transformation by Buffering

Implements:

TransformationHandler, **instance** th .

Uses:

JobHandler, **instance** jh .

upon event $\langle th, Init \rangle$ **do**

$top := 1;$

$bottom := 1;$

$handling := \text{FALSE};$

$buffer := [\perp]^M;$

upon event $\langle th, Submit | job \rangle$ **do**
if $bottom + M = top$ **then**
 trigger $\langle th, Error | job \rangle$;
else
 $buffer[top \bmod M + 1] := job;$
 $top := top + 1;$
 trigger $\langle th, Confirm | job \rangle$;

upon $bottom < top \wedge handling = \text{FALSE}$ **do**

$job := buffer[bottom \bmod M + 1];$

$bottom := bottom + 1;$

$handling := \text{TRUE};$

trigger $\langle jh, Submit | job \rangle$;

upon event $\langle jh, Confirm | job \rangle$ **do**

$handling := \text{FALSE};$

Joke time!

“Distributed systems would be easy —

“Distributed systems would be easy –
if time existed.”

Protocol vs Distributed Algorithm

To Provide Service

- A module at a process
 - Executes one or more **rounds** of message exchanges with
 - Peer modules at other/remote processes

Protocol is ...

- The behavior of each peer characterized by
 - the set of messages each peer is capable of producing and accepting,
 - the format of these messages, and
 - the legal sequences of messages

Protocol's Purpose is to ...

- Ensure the execution of some **distributed algorithm**.
- Concurrent execution of different sequences of steps
 - that ensure the provision of the desired service

Classes of Algorithmic Solutions

- Book provides insight into how
 - **Failure assumptions**, the **environment**, the system parameters, and
 - other design choices *affect the algorithm design ...*
- Several different classes of algorithmic solutions to the
 - distributed programming abstractions

Classes of Algorithmic Solutions

1. **Fail-stop**: processes can *fail by crashing* but the crashes can be *reliably detected*
2. **Fail-silent**: process crashes can *never be reliably detected*
3. **Fail-noisy**: crashes can be detected, but not always in an accurate manner
4. **Fail-recovery**: processes can crash and later recover and participate in the algorithm
5. **Fail-arbitrary**: processes can deviate arbitrarily from the protocol specification and act in malicious, adversarial ways
6. **Randomized** algorithms: processes may make probabilistic choices

Questions?